amazon

HULT
INTERNATIONAL
BUSINESS SCHOOL

Table of Contents

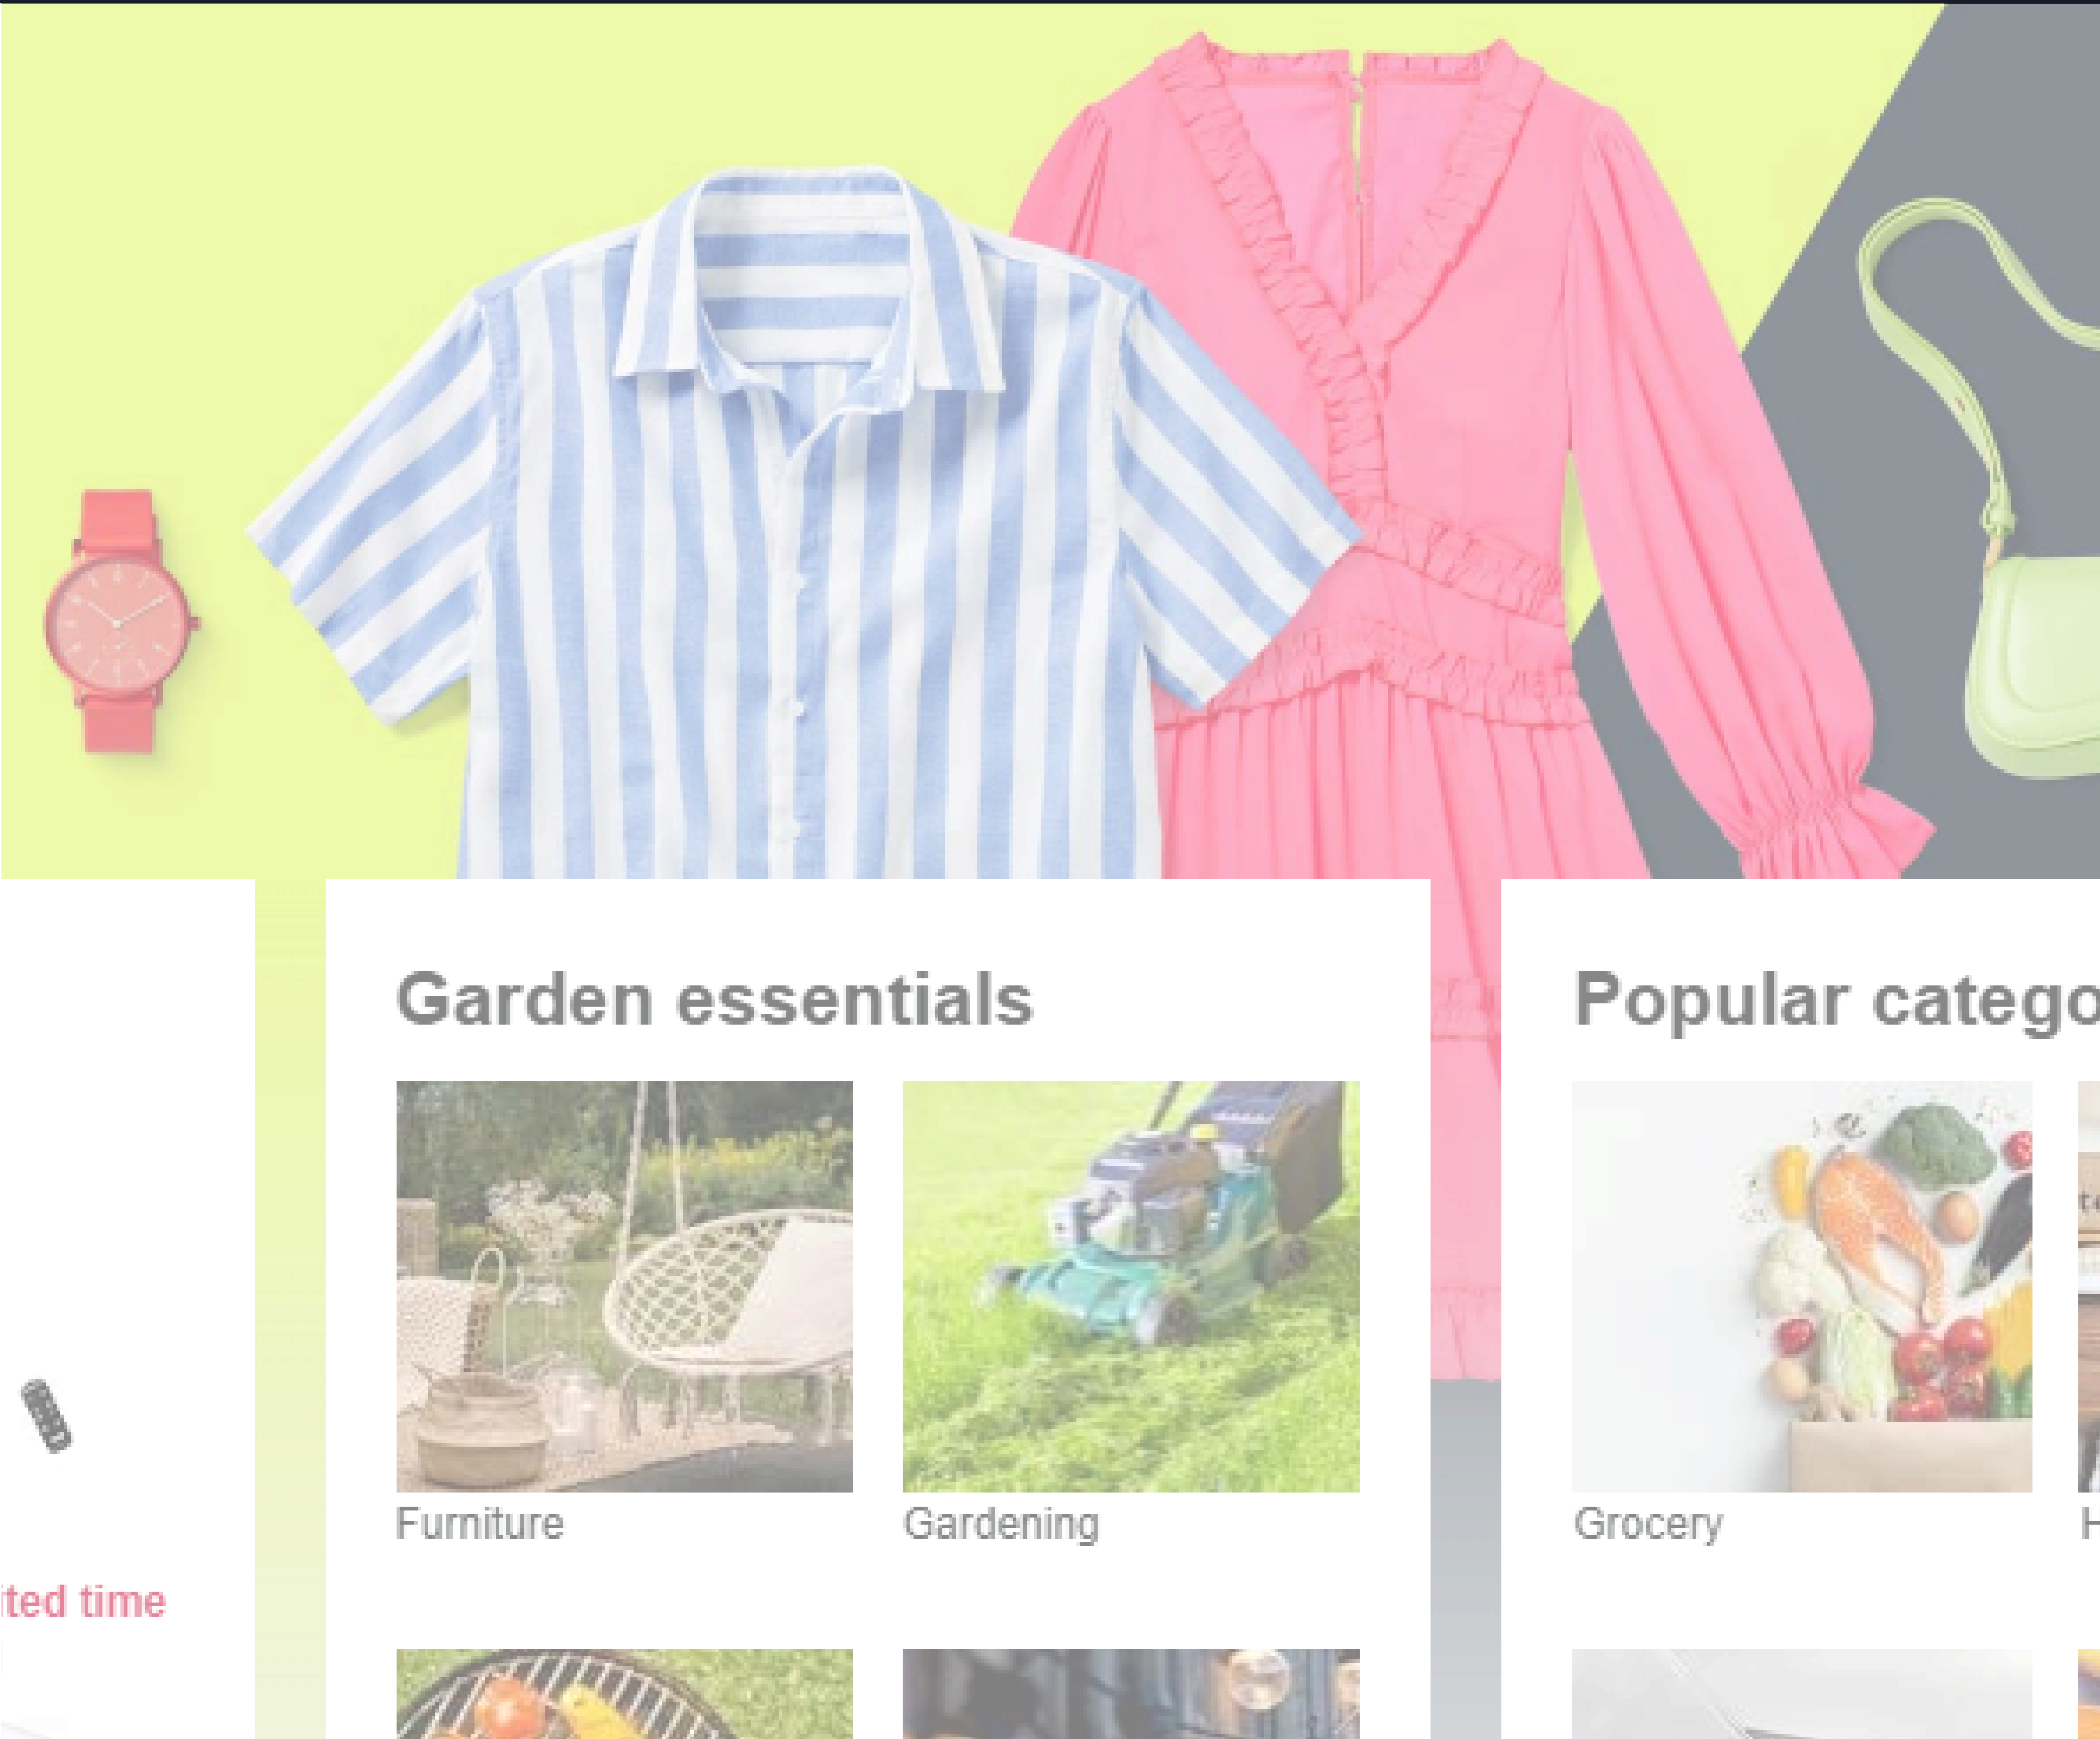

amazon Delivering to HULT INTERNATIONAL BUSINESS SCHOOL

All Introduction

0 Team 3 Basket

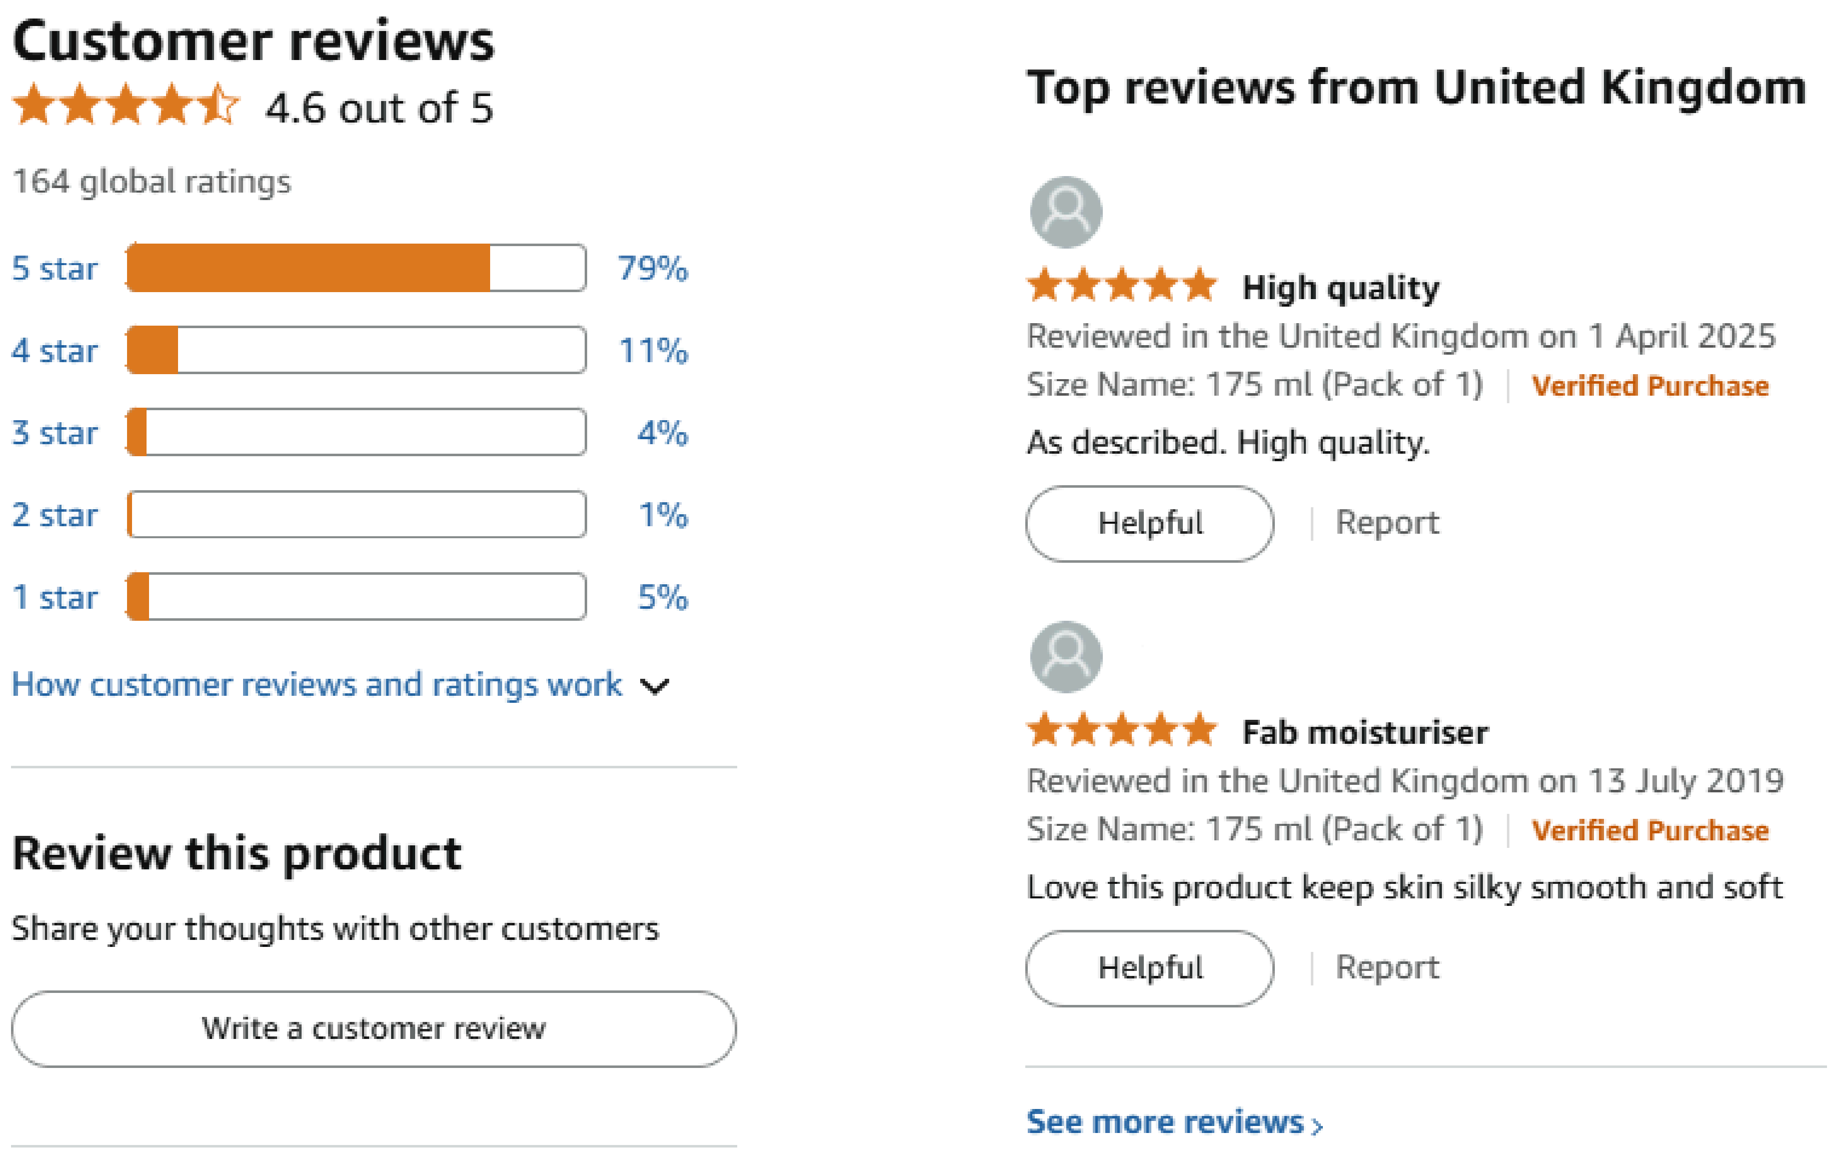

Business Problem:
Consumers rely heavily on product reviews when making purchasing decisions, yet these reviews can vary widely in quality, sentiment, and usefulness. Businesses often lack insights into:

- What makes a review helpful or influential?
- How customer sentiment shifts over time and across products?
- Which aspects of a product drive positive or negative feedback?

Objective:
To analyze Amazon food product reviews to uncover trends in review helpfulness, sentiment dynamics, and thematic topics, helping businesses enhance customer satisfaction, product development, and marketing strategies.

Delivering to HULT INTERNATIONAL BUSINESS SCHOOL

All

Dataset Overview

0 Team 3 Basket

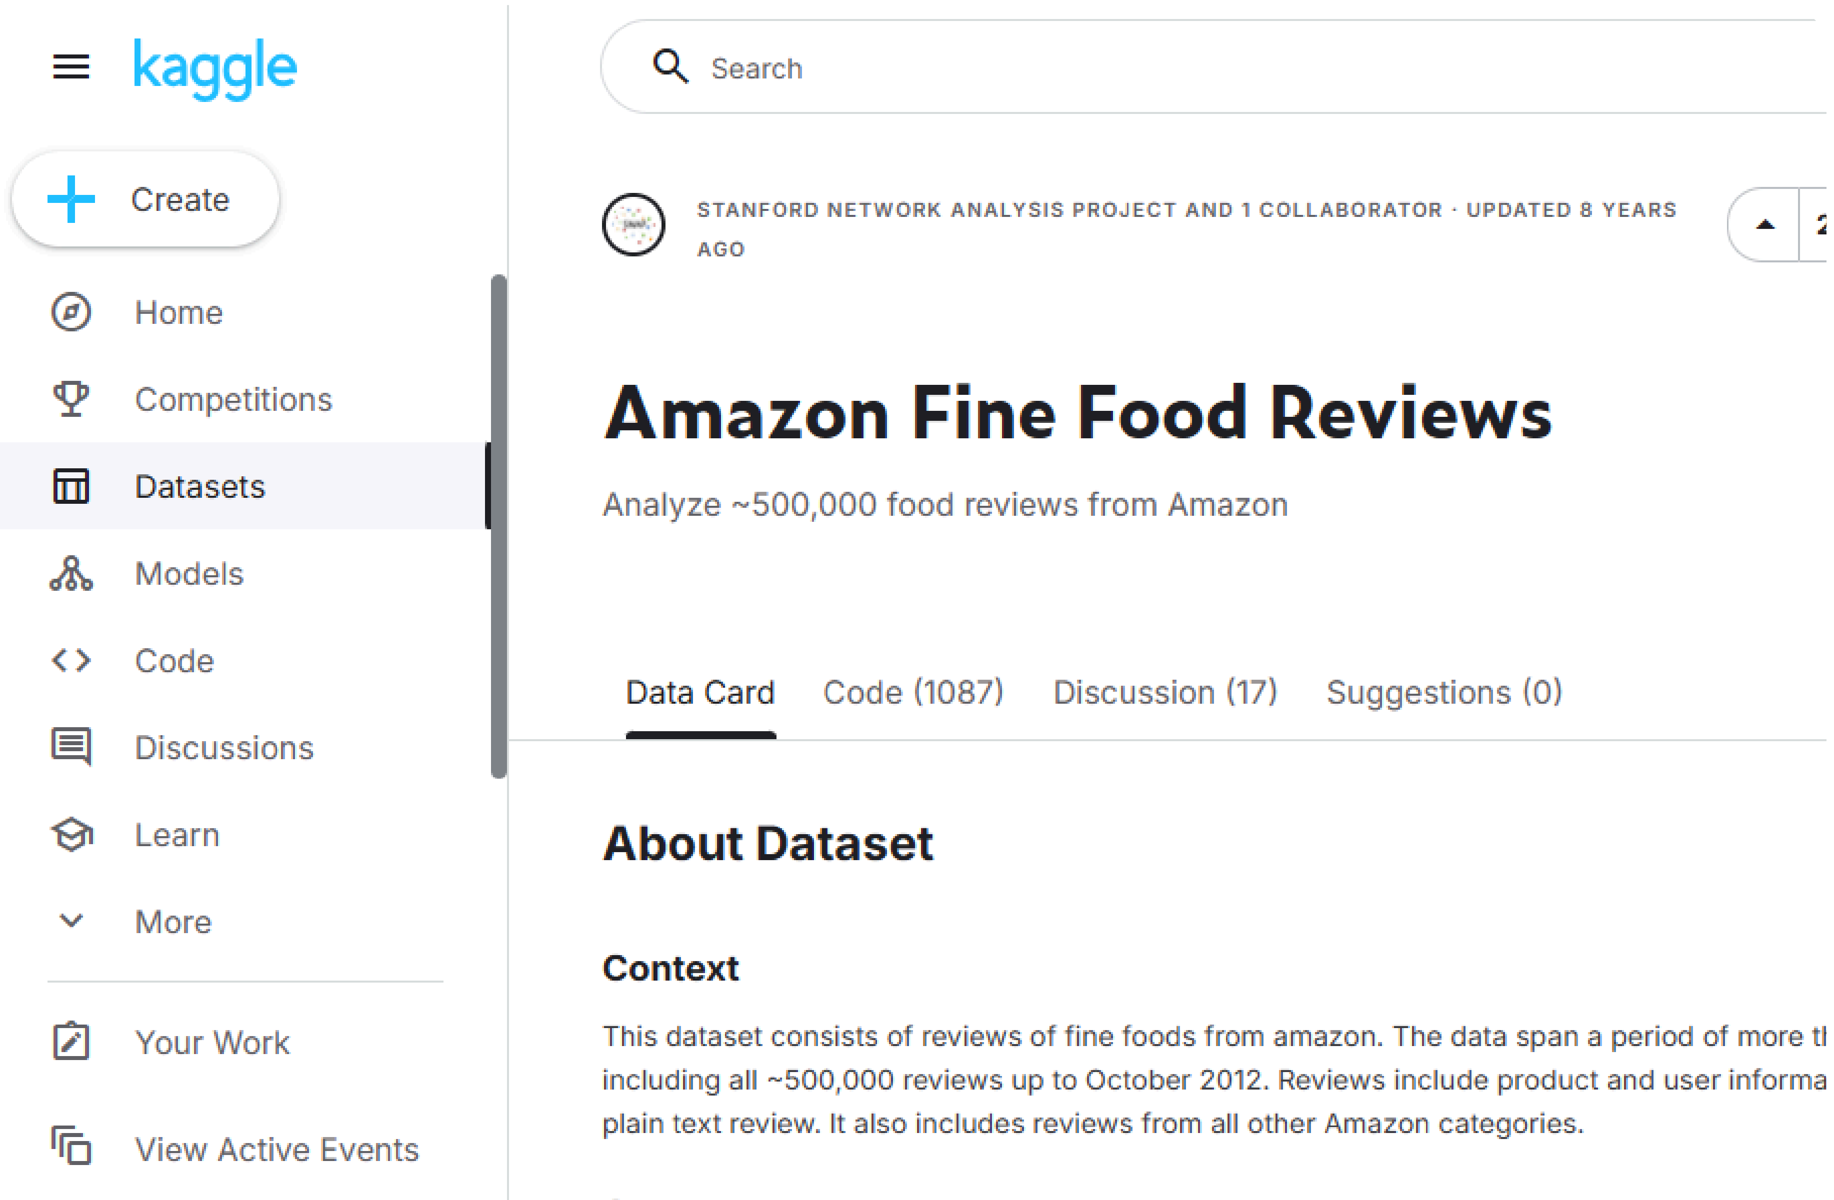

https://www.kaggle.com/datasets/snap/amazon-fine-food-reviews?resource=download

- **Source:** Amazon Fine Food Reviews dataset (Kaggle)
- **Records:** ~568,000 reviews of food-related products
- **Time Span:** 1999 to 2012
- **Key Fields:**
  - ProductId: Unique ID of the product
  - UserId & ProfileName: Reviewer identifiers
  - Score: Star rating (1–5)
  - Time: Review timestamp
  - Summary & Text: Review content
  - HelpfulnessNumerator / Denominator: Peer-rated helpfulness

Data Pre-Processing

HULT
INTERNATIONAL
BUSINESS SCHOOL

Data Pre-Processing

- Converted all review text to lowercase
- Removed punctuation, non-alphabetic characters, and English stopwords using NLTK
- Created a new column: Cleaned_Text using a custom clean_text() function
- Applied tokenization and text standardization
- Trimmed dataset to the first 600,000 records to optimize performance
- Output: Clean, consistent text for TF-IDF, sentiment analysis, and topic modeling

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary	Text
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1	5	1303862400	Good Quality Dog Food	I have bought several of the Vitality canned d...
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0	1	1346976000	Not as Advertised	Product arrived labeled as Jumbo Salted Peanut...
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1	4	1219017600	"Delight" says it all	This is a confection that has been around a fe...
3	4	B000UA0QIQ	A395BORC6FGVXV	Karl	3	3	2	1307923200	Cough Medicine	If you are looking for the secret ingredient i...
4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham "M. Wassir"	0	0	5	1350777600	Great taffy	Great taffy at a great price. There was a wid...

Feature Engineering

- TF-IDF Vectorization
  - Applied to Cleaned_Text
  - Limited to top 5,000 terms for relevance and efficiency
- Sentiment Label Creation
  - Based on Score:
    - Positive (4–5)
    - Neutral (3)
    - Negative (1–2)
- Helpfulness_Ratio
  - Formula: HelpfulnessNumerator / HelpfulnessDenominator
  - Capped at 1.0 to handle anomalies
- ReviewLength & SummaryLength
  - Counted words in full review and summary respectively
- ReviewAgeDays
  - Calculated time since review using the current date
- All features supported:
  - Sentiment classification
  - Topic modeling
  - Trend analysis
  - User behavior profiling

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary	Text
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1	5	1303862400	Good Quality Dog Food	I have bought several of the Vitality canned d...
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0	1	1346976000	Not as Advertised	Product arrived labeled as Jumbo Salted Peanut...
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1	4	1219017600	"Delight" says it all	This is a confection that has been around a fe...
3	4	B000UA0QIQ	A395BORC6FGVXV	Karl	3	3	2	1307923200	Cough Medicine	If you are looking for the secret ingredient i...
4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham "M. Wassir"	0	0	5	1350777600	Great taffy	Great taffy at a great price. There was a wid...

Exploratory Data Analysis: Foundational Distribution

Foundational Distribution: Score Distribution

Overwhelmingly Positive Ratings Dominate Review Scores

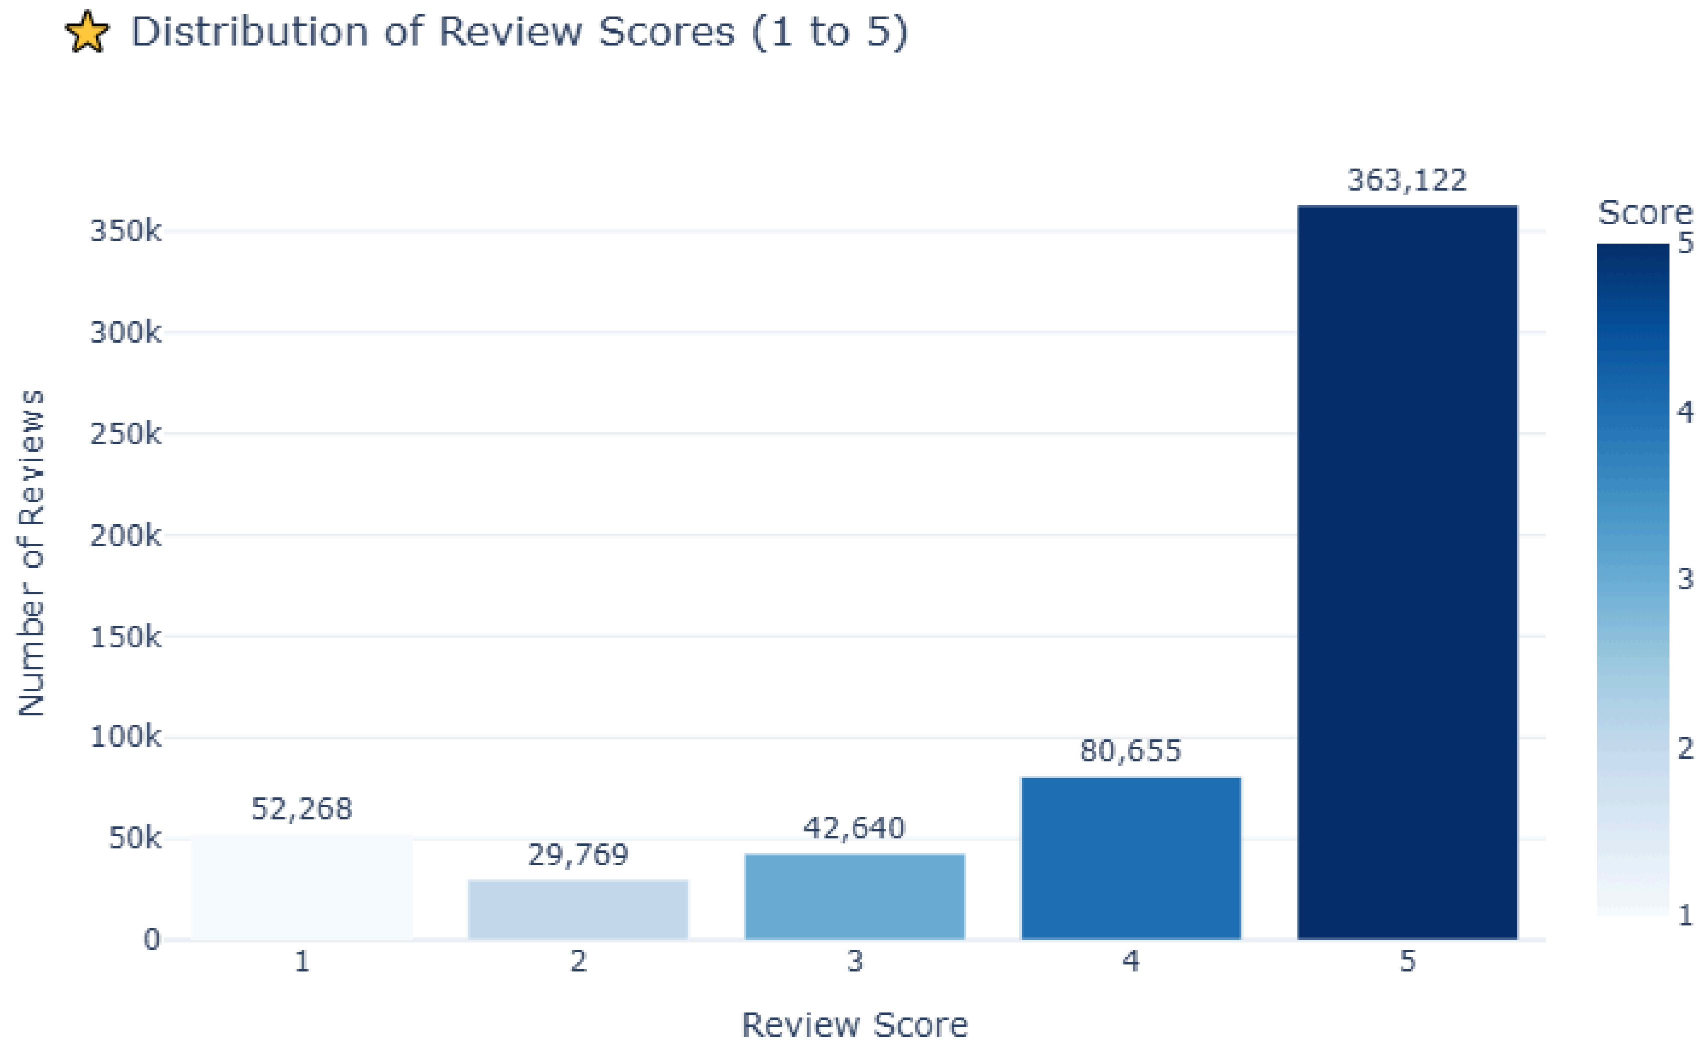

- Score distribution is heavily skewed toward positive ratings
  - Majority of reviews are 5-star, followed by 4-star
- Low occurrence of 1- and 2-star ratings
  - Indicates fewer customers leave negative feedback
- Neutral (3-star) reviews are also underrepresented
  - Reflects a tendency toward polarized opinions
- Implications for modeling:
  - Class imbalance may affect sentiment classifier performance
- Business insight:
  - Relying solely on score averages may overstate satisfaction
  - Textual and helpfulness analysis needed for richer insights

Foundational Distribution: Sentiment Distribution

Customer Reviews Skew Heavily Toward Positivity

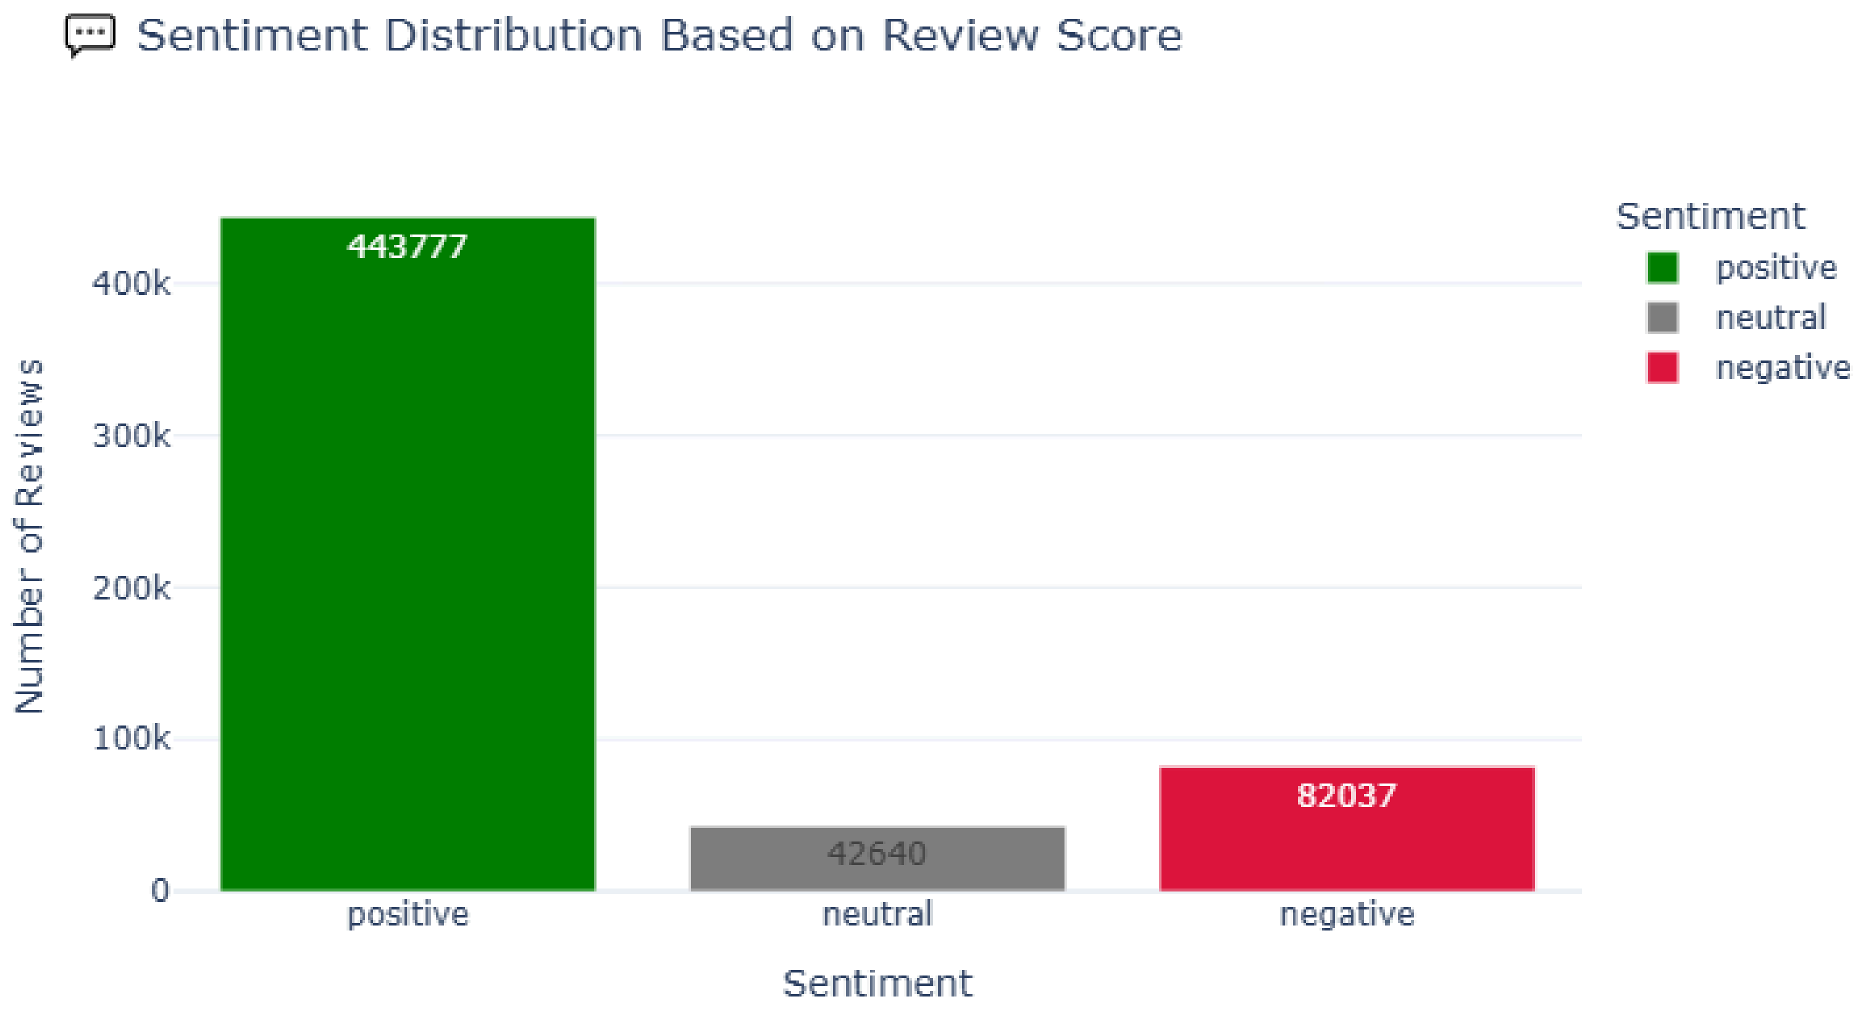

- 78% of reviews classified as positive
  - Aligns with the skew seen in star ratings
- Negative (14%) and neutral (7.5%) sentiments are underrepresented
  - Highlights polarization in customer feedback
- Possible reasons:
  - Review bias – users with strong opinions more likely to post
  - Genuine satisfaction, but may mask critical feedback
- Implications for modeling:
  - Imbalanced sentiment classes affect classifier accuracy
- Business takeaway:
  - Positive sentiment is promising, but outliers may hold key insights
  - Text analysis and helpfulness ratings provide deeper quality signals

Foundational Distribution: Helpfulness Ratio Distribution

Review Helpfulness Is All or Nothing

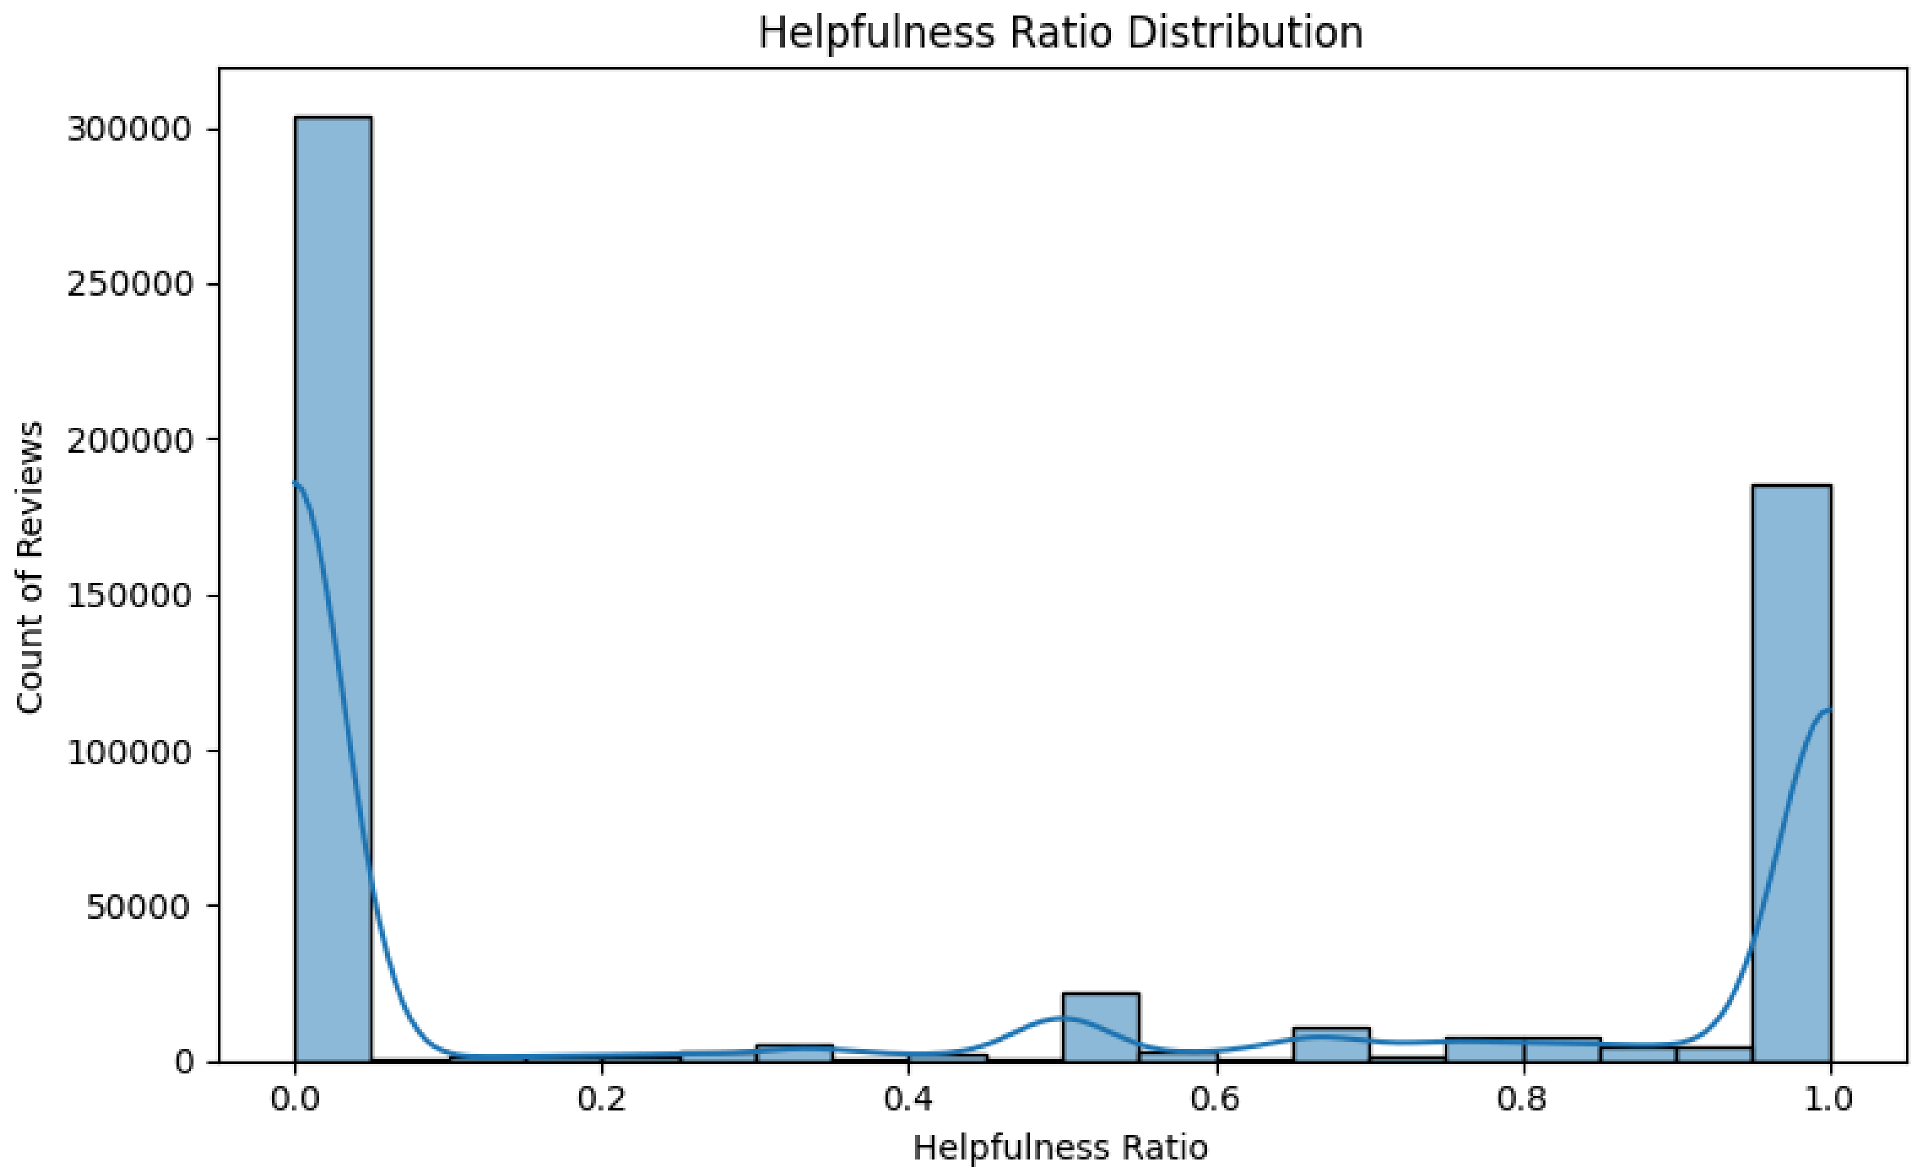

- Helpfulness ratio shows a bimodal distribution:
  - Many reviews have a ratio of 0.0 (no helpful votes)
  - Others have a perfect ratio of 1.0 (all votes marked helpful)
- Indicates reviews are either highly valued or largely ignored
- Influenced by factors like:
  - Length, clarity, and timing of the review
- Short or vague reviews often get no engagement
- Well-written, detailed feedback more likely to be rated helpful
- Implication: Helpfulness ratio can act as a proxy for review quality
- But beware: many reviews go unrated, which may skew perception

Exploratory Data Analysis: Review Content Analysis

Review Content Analysis: Distribution of Review and Summary Lengths

Majority of Reviews Are Longer Than Summaries, But Still Brief Overall

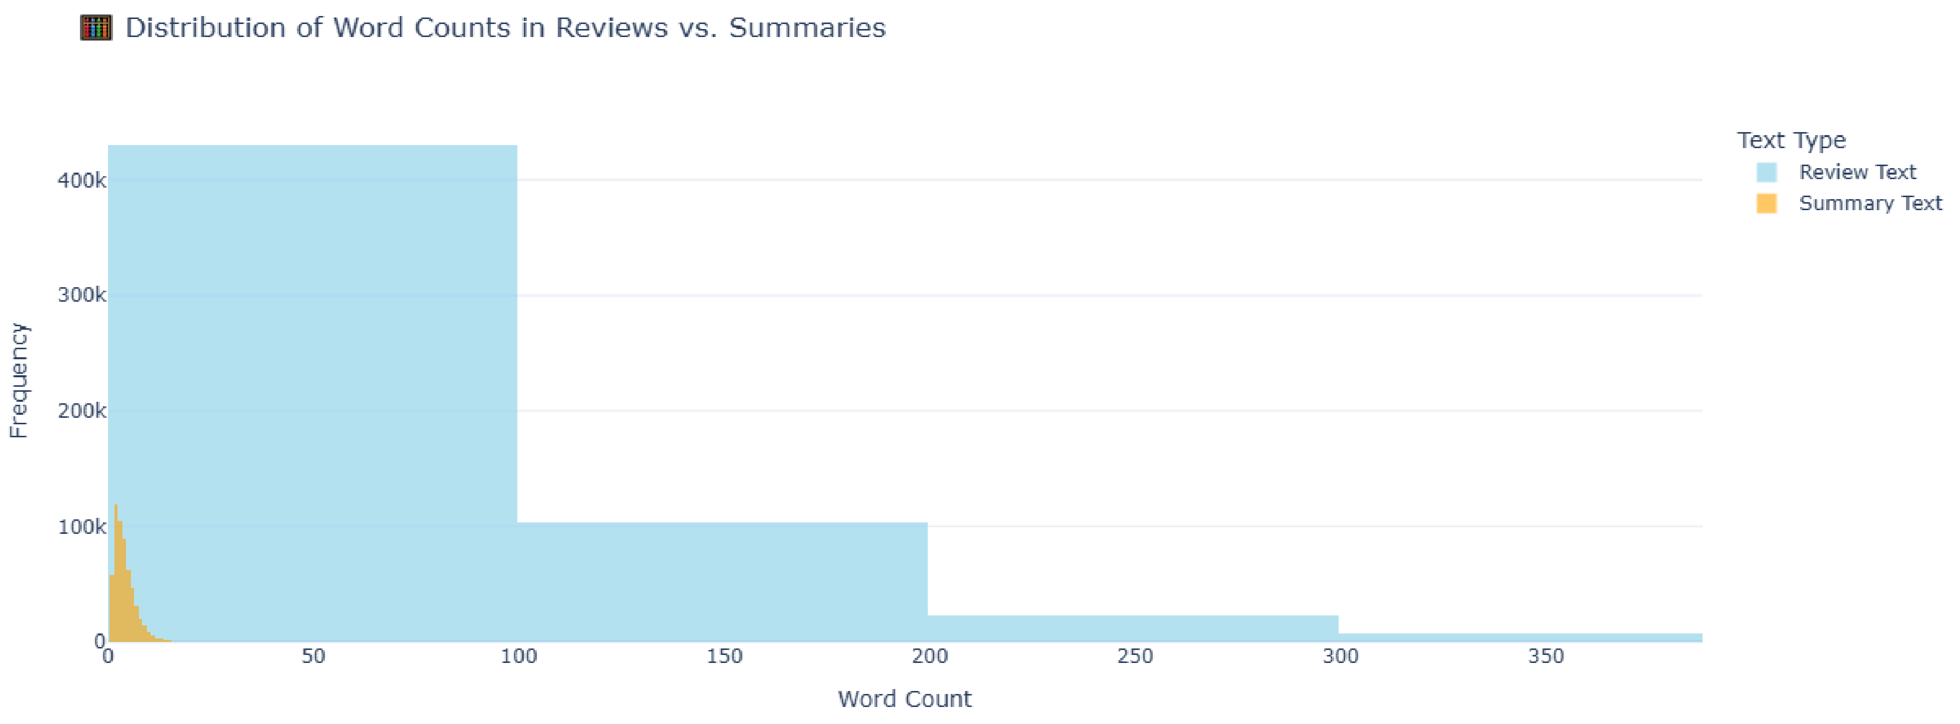

- Reviews are longer, typically 20–100 words; summaries are much shorter (often <20 words).
- Both are right-skewed, showing most users write brief content.
- Summaries offer quick sentiment cues, but reviews provide richer insights.
- Short reviews may limit sentiment depth, highlighting the need for text analysis techniques.
- Encouraging detailed reviews can enhance sentiment accuracy and business insights.

Review Content Analysis: Review Length vs Helpfulness

Longer Reviews, Higher Helpfulness?

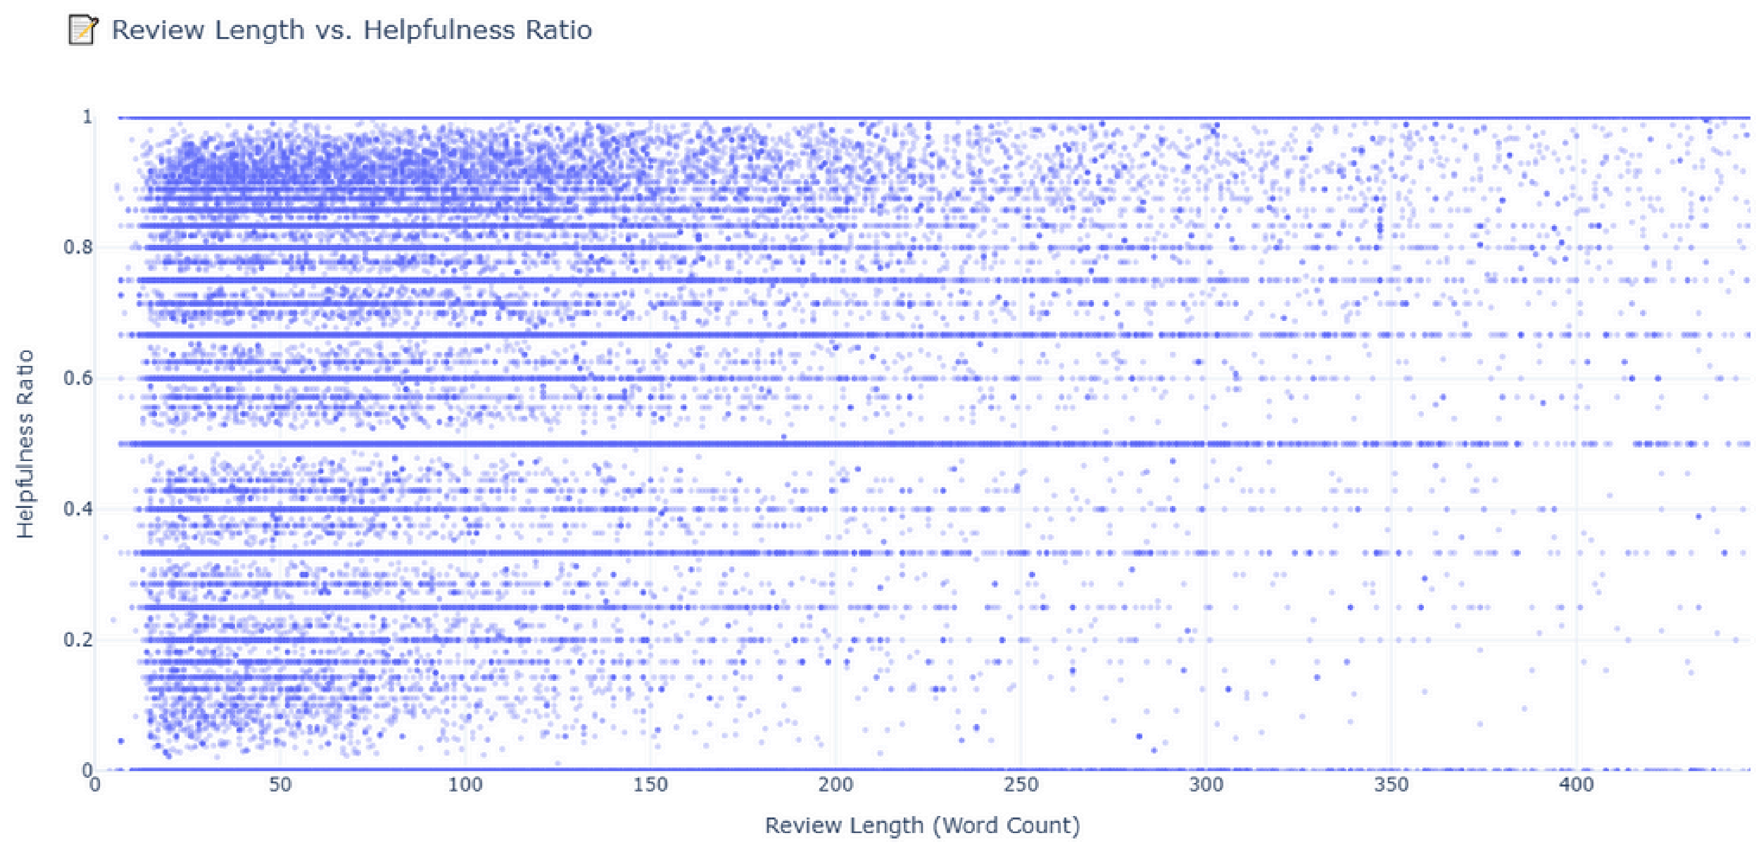

- Positive correlation observed between review length and helpfulness ratio—longer reviews tend to be rated as more helpful.
- Reviews exceeding 100+ words show a higher density near a helpfulness ratio of 1.0, suggesting perceived value in detailed feedback.
- Shorter reviews cluster around low helpfulness scores, possibly due to lack of context or clarity.
- Despite variability, a trend emerges: quality often lies in depth, benefiting product credibility and future buyers.
- For businesses, encouraging informative reviews may improve conversion influence and trust-building with potential customers.

Exploratory Data Analysis: User and Product Activity

User and Product Activity: Most Reviewed Users

Power Users Drive Review Volume

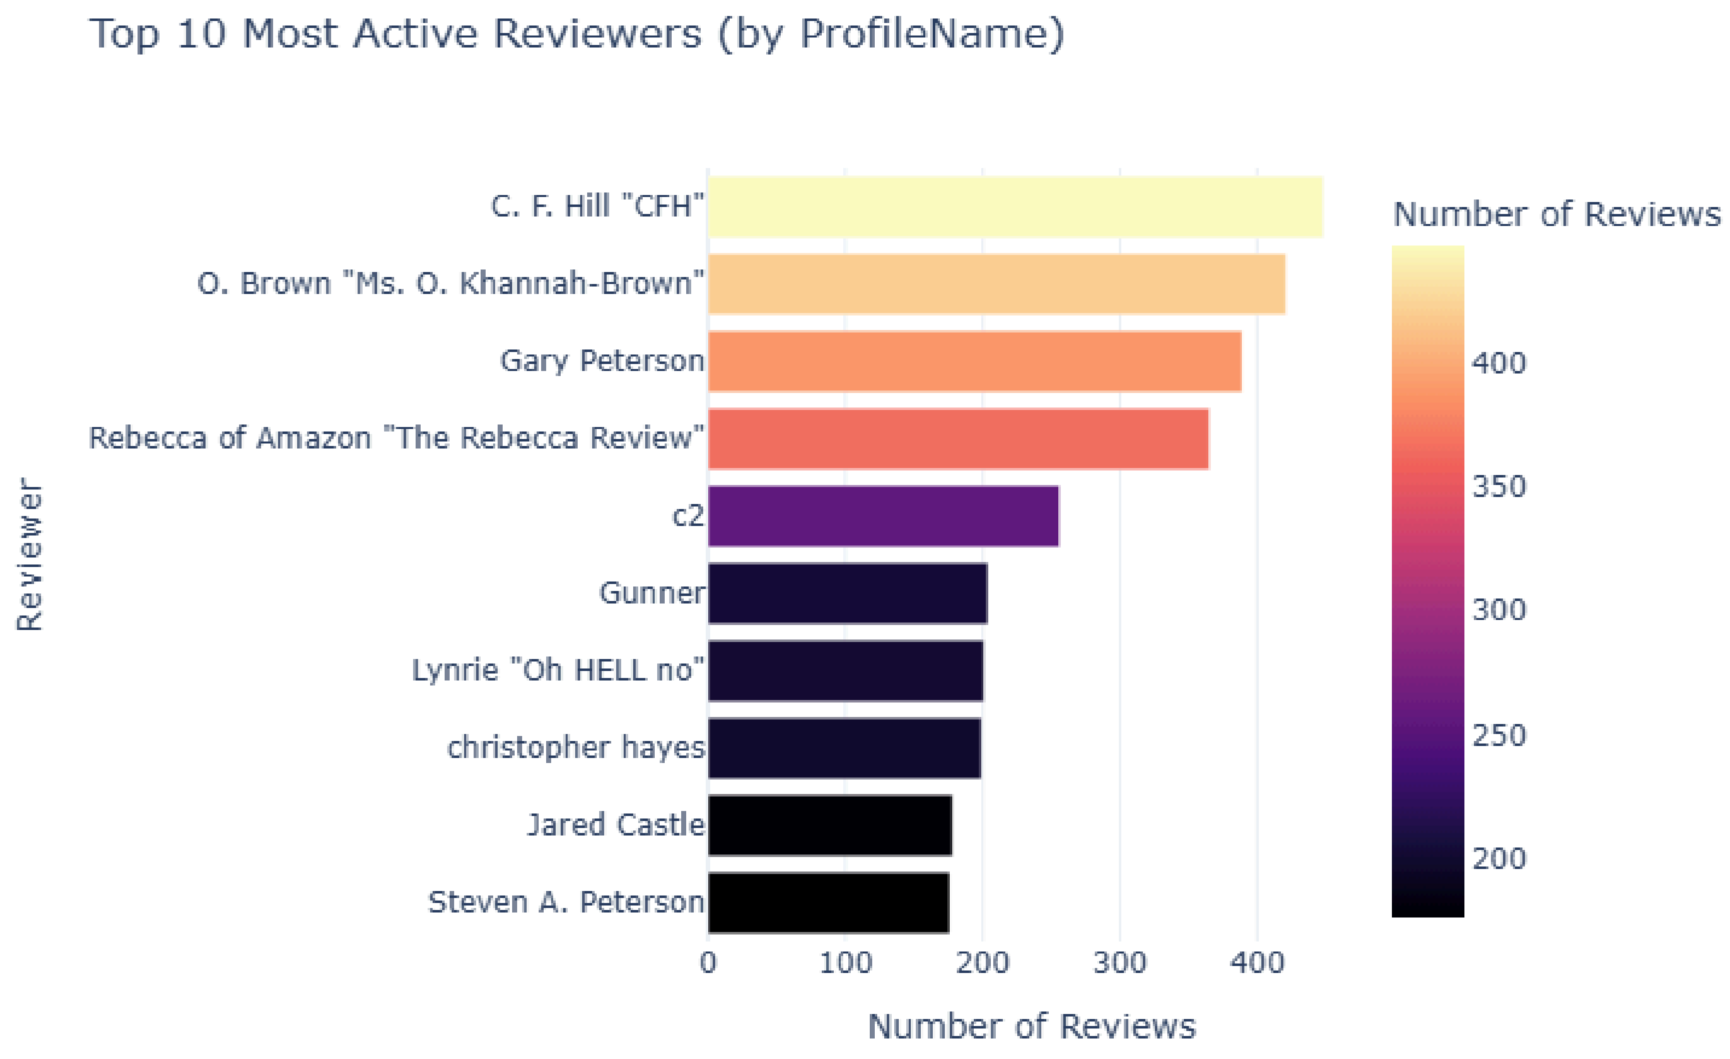

- The top 10 most active reviewers each contributed between 180 and 450 reviews, indicating a small segment of users account for a substantial portion of feedback.
- C. F. Hill ("CFH") led the board with the highest number of reviews, suggesting potential influence on product visibility and perception.
- Many top contributors use branded or stylized names (e.g., “The Rebecca Review”), hinting at semi-professional or enthusiast reviewers.
- Businesses could leverage engagement with prolific reviewers for early product testing or outreach campaigns.
- These reviewers represent valuable community voices, possibly shaping purchasing decisions more than casual users.

User and Product Activity: Most Reviewed Products

Most Reviewed Products Dominate User Attention

```
Top 10 Most Reviewed Products:
ReviewCount  count
 B007JFMH8M    913
 B0026RQTGE    632
 B002QWHJOU    632
 B002QWP89S    632
 B002QWP8H0    632
 B003B3OOPA    623
 B001EO5Q64    567
 B000VK8AVK    564
 B0026KNQSA    564
 B007M83302    564
```

- The top product, B007JFMH8M, received 913 reviews, significantly more than the next most reviewed items.
- A cluster of products (e.g., B0026RQ series) shared identical review counts (632 each), suggesting coordinated campaigns, bundled listings, or popular variations.
- High review volume reflects strong customer engagement and potentially high sales or visibility on the platform.
- These products are ideal candidates for targeted sentiment and quality analysis, as the volume offers rich data for trend extraction.
- Businesses can focus on these high-volume items to optimize messaging, improve quality, and monitor reputation closely.

User and Product Activity: Review Length by Review Score

Review Length Varies with Score Intensity

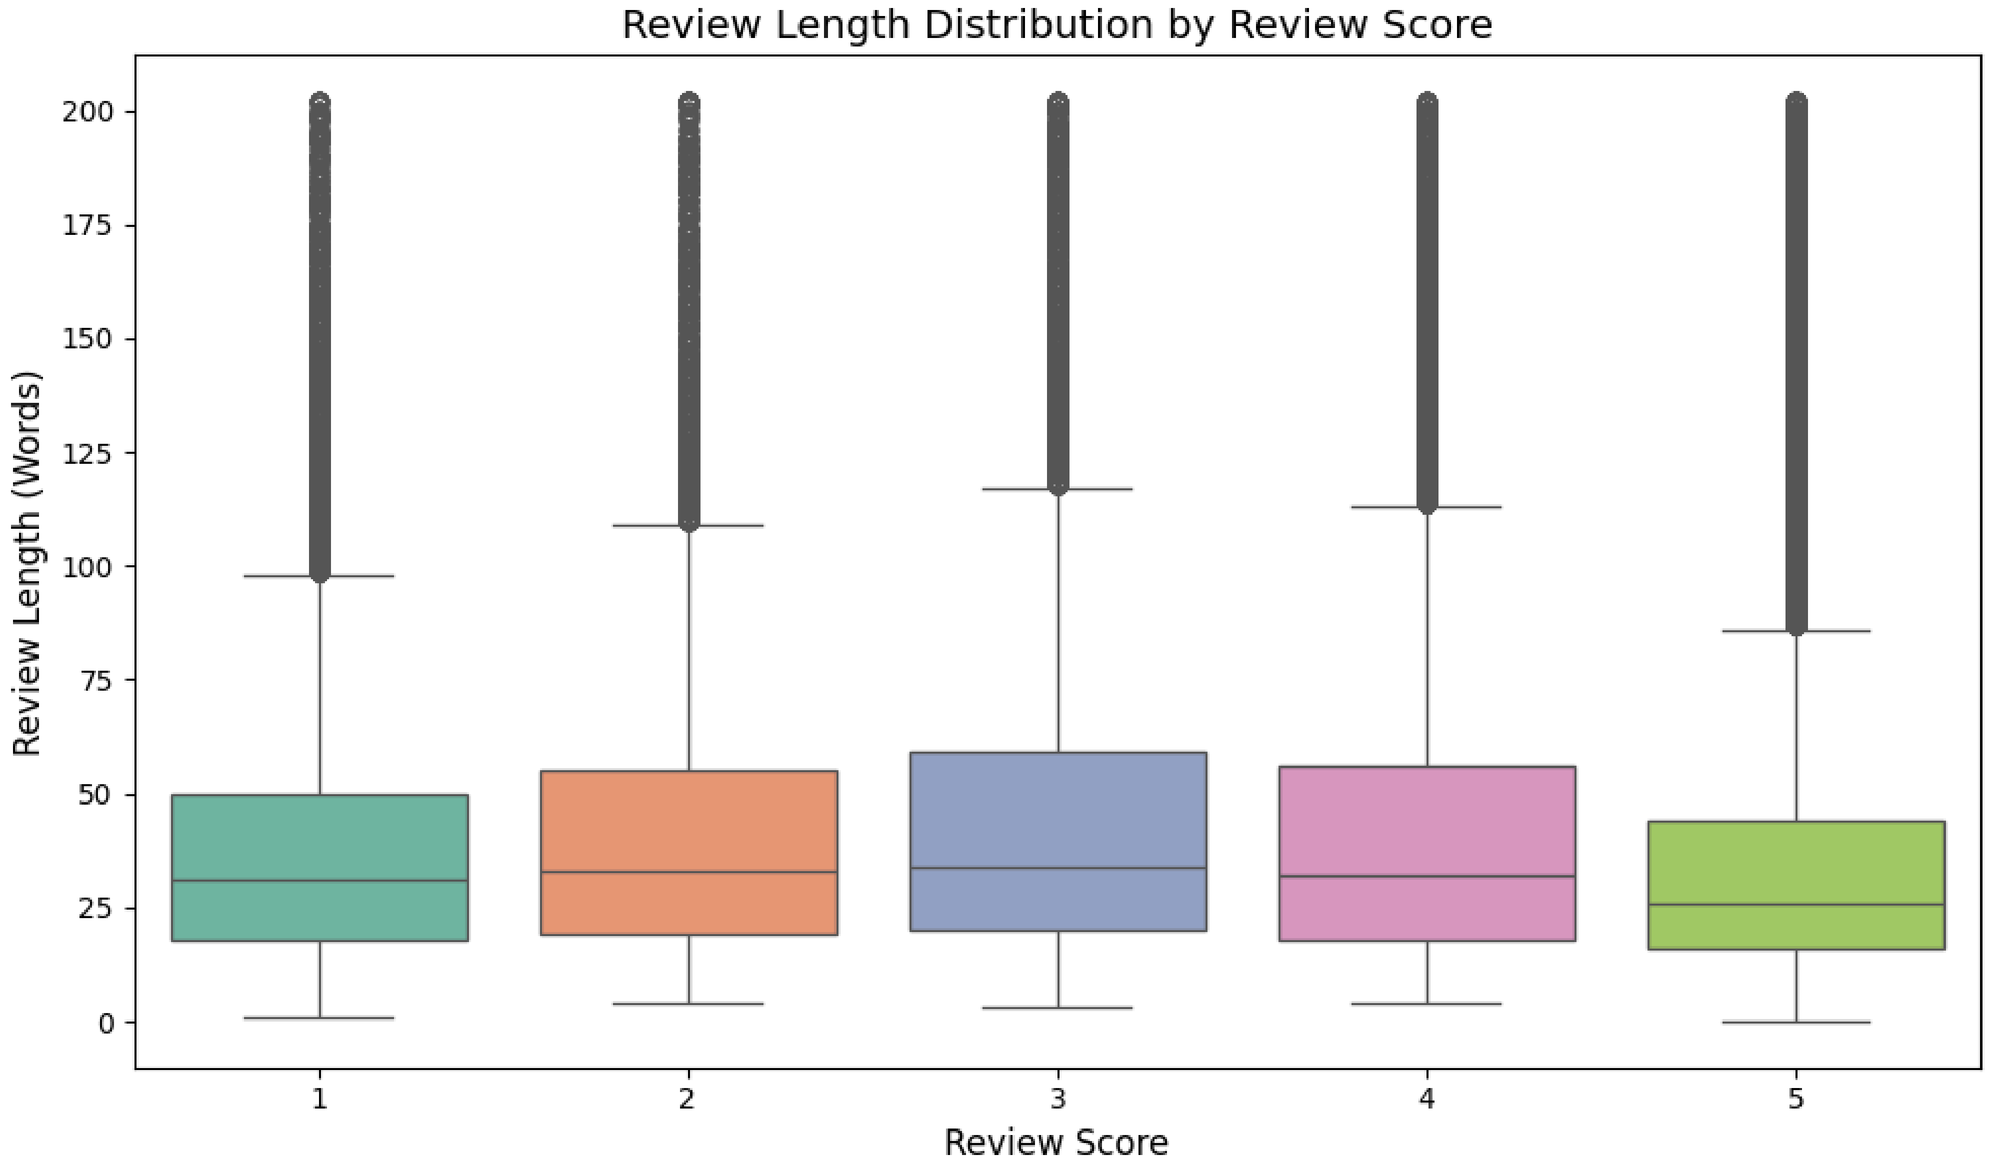

- Reviews with extreme scores (1 or 5) tend to be shorter, likely reflecting quick expressions of strong opinions.
- Moderate scores (2–4) show longer median lengths, suggesting more detailed reasoning or balanced perspectives.
- Outliers with over 200 words exist across all score levels, indicating occasional in-depth reviews regardless of sentiment.
- This variation helps businesses identify that lengthy reviews may signal nuanced or constructive feedback, while shorter ones may highlight intense praise or dissatisfaction.
- Prioritizing analysis of mid-range, longer reviews could unlock richer product improvement insights.

User and Product Activity: Average Score per User

Most Users Are Generous Raters

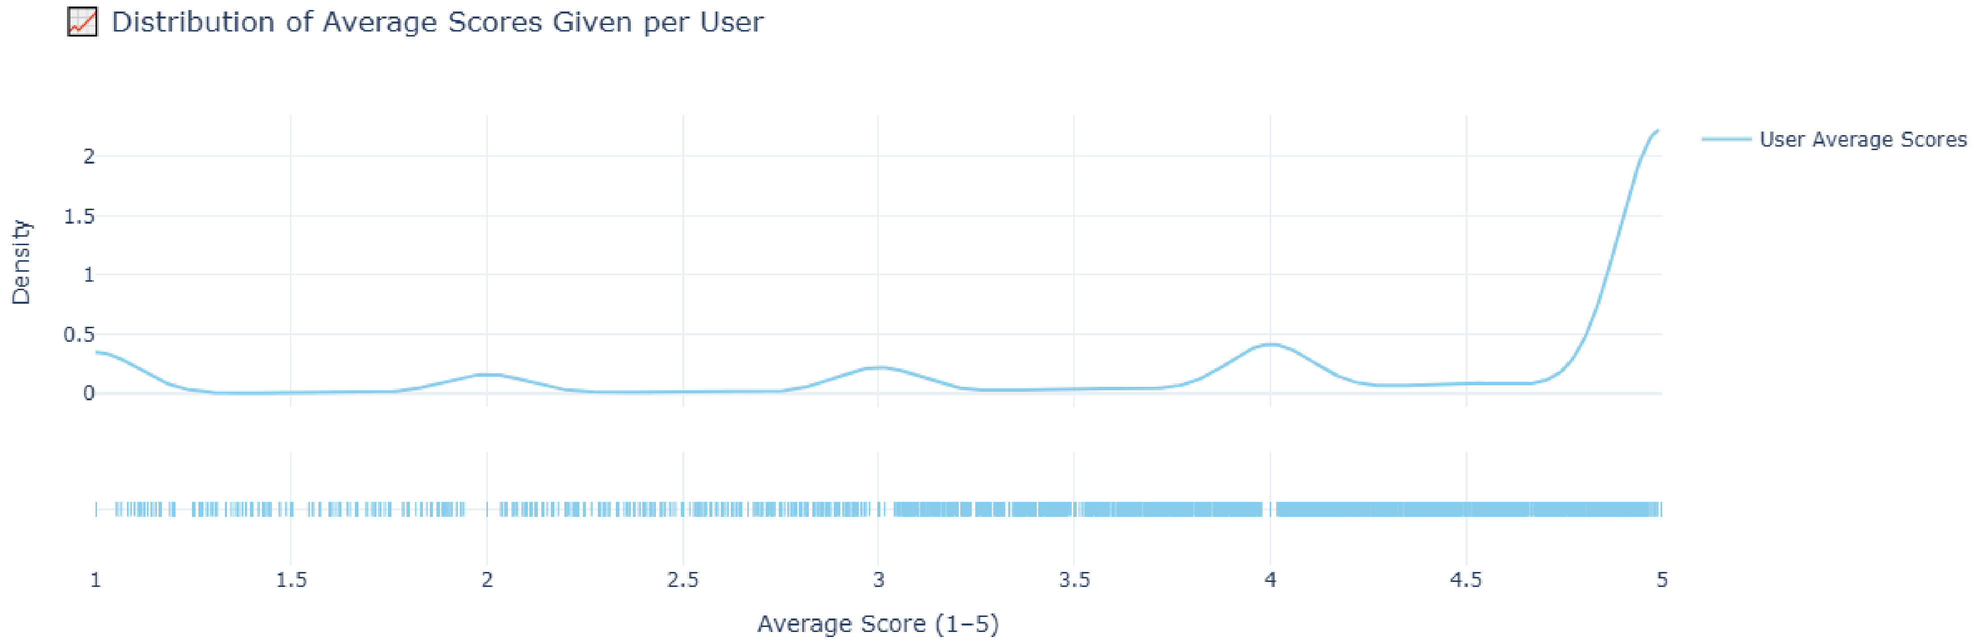

- The density plot shows a heavy concentration of users giving average scores close to 5, confirming a strong positive bias in user behavior.
- A smaller number of users provide lower average scores, with minor density peaks near 2, 3, and 4, suggesting a few balanced or critical reviewers.
- This skewness implies that most customers only review when they're highly satisfied, which may inflate a product's perceived quality.
- Businesses should monitor consistent low scorers for potential product issues and weigh positive reviews cautiously, especially when they dominate.
- Understanding these patterns can inform review credibility scoring systems and help identify loyal brand advocates vs. critical evaluators.

Exploratory Data Analysis: Text and Sentiment Features

HULT
INTERNATIONAL
BUSINESS SCHOOL

Text and Sentiment Features: Most Frequent Words

What Customers Talk About Most

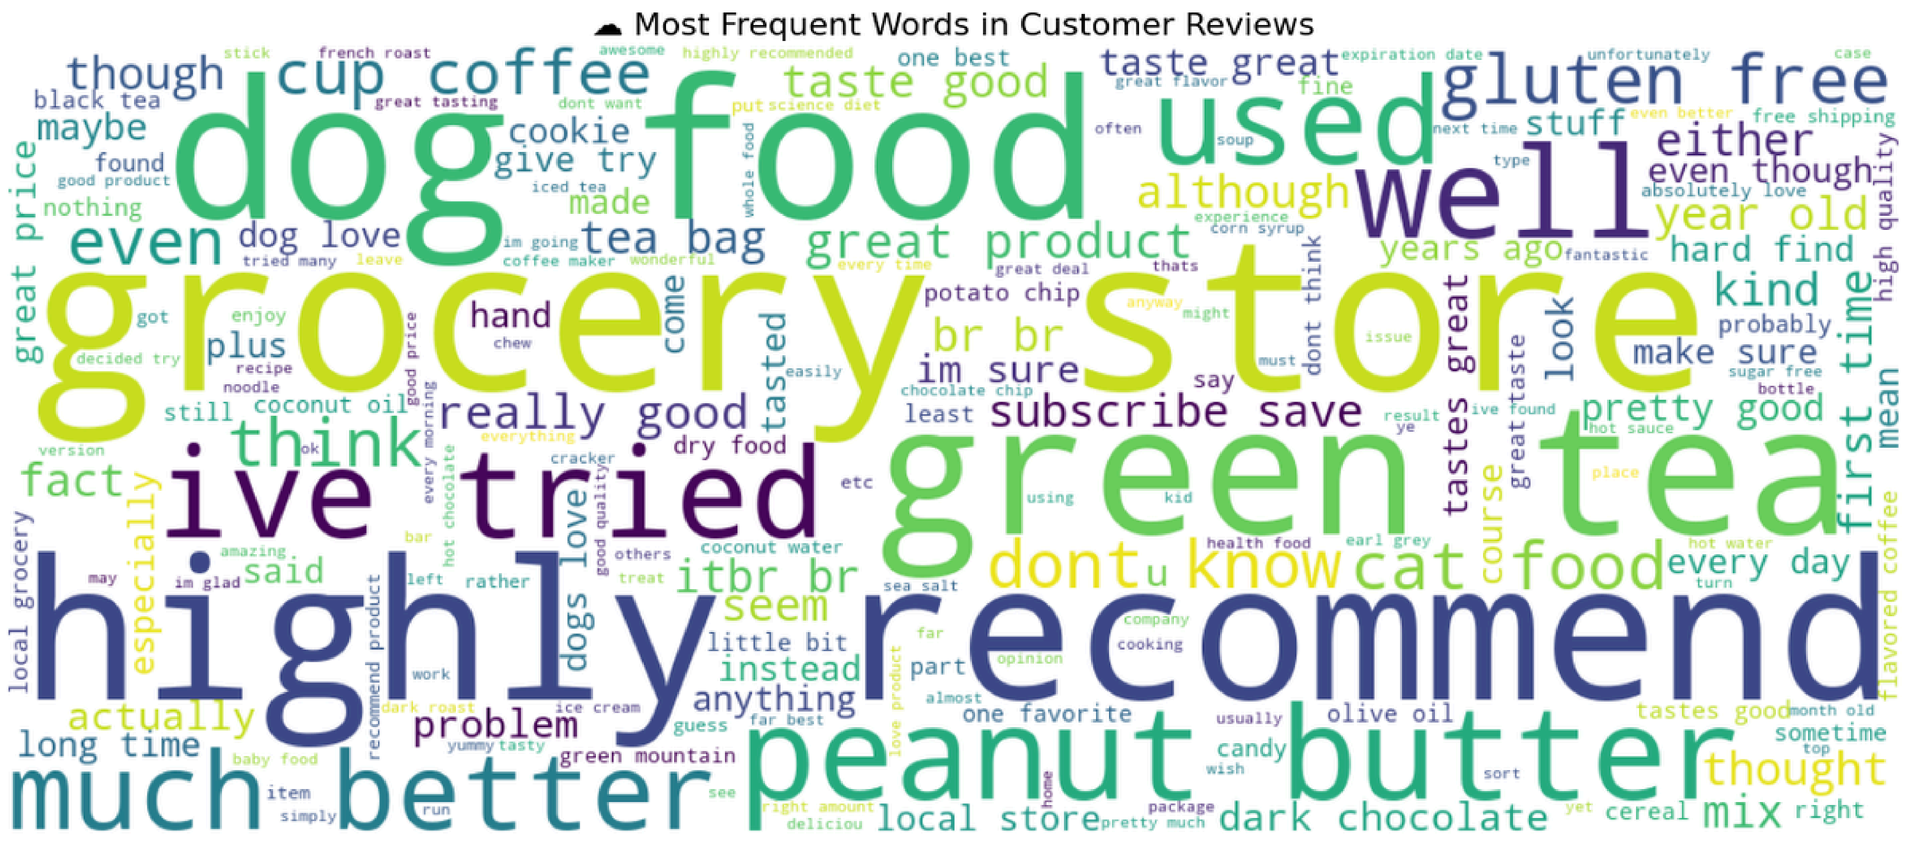

- Frequent keywords like “recommend,” “store,” “grocery,” “highly,” “dog food,” “green tea,” and “peanut butter” dominate the word cloud.
- Words such as “subscribe,” “save,” “taste,” “love,” and “product” suggest strong satisfaction and recurring purchase behavior.
- Presence of product-specific terms (e.g., "tea bag," "coconut oil") indicates high engagement with consumables and pet-related items.
- Positive adjectives like “great,” “good,” “much better,” “really” reinforce the earlier findings of sentiment skew toward positive feedback.
- This confirms customers are not only reviewing favorably, but also frequently recommending and endorsing product value, which companies can leverage in marketing messaging.

Text and Sentiment Features: Helpfulness by Score

Positive Reviews Are More Helpful

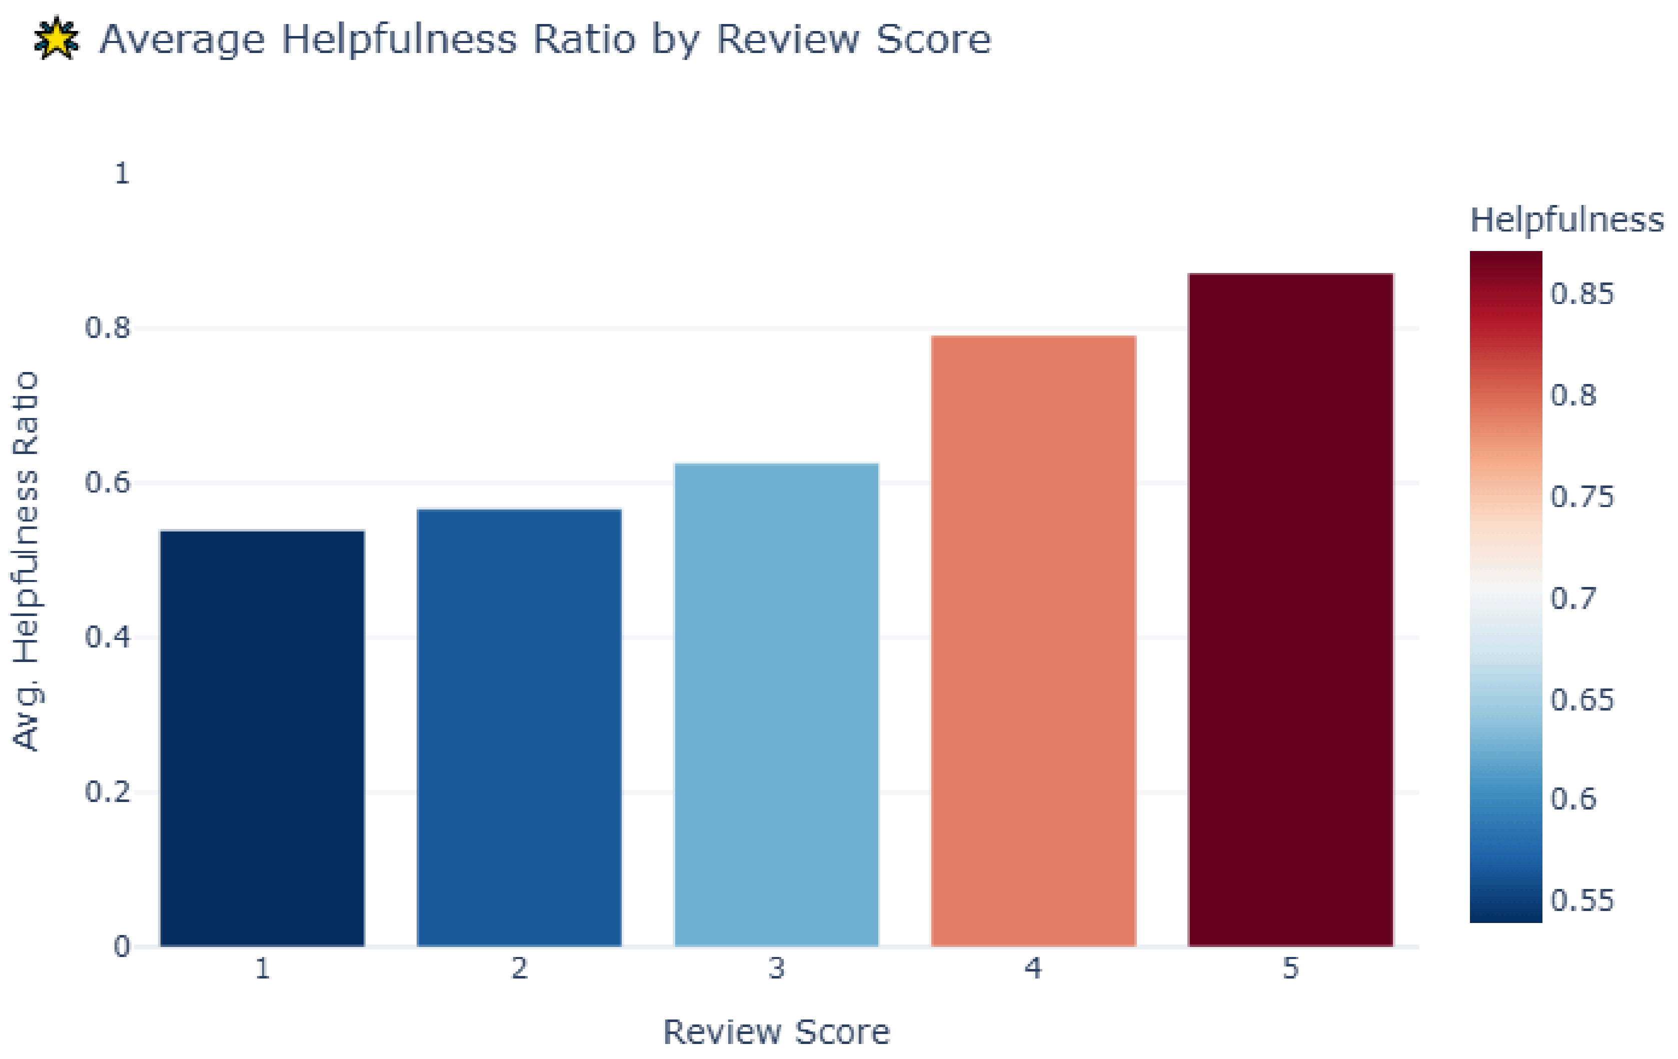

- Reviews with higher star ratings (4 and 5) tend to have higher average helpfulness ratios, suggesting that positively framed feedback resonates more with readers.
- The helpfulness ratio rises steadily with increasing review scores, peaking at nearly 0.88 for 5-star reviews.
- Lower-scoring reviews (1–2 stars) are perceived as less helpful, possibly due to brevity, emotional tone, or lack of constructive detail.
- From a business lens, this insight highlights that well-articulated positive reviews can amplify product credibility and may influence purchasing decisions.
- Encouraging detailed reviews—regardless of sentiment—could enhance user trust and improve content value across all score levels.

Delivering to HULT INTERNATIONAL BUSINESS SCHOOL

All Exploratory Data Analysis

0 Team 3 Basket

Text and Sentiment Features: Sentiment Over Time

Surge in Positive Sentiment Over Time

- Positive reviews skyrocketed after 2006, peaking at over 150,000 in 2012, indicating growing customer satisfaction or review participation.
- Negative and neutral sentiments also increased, but at a much slower rate, with negative reviews reaching around 32,000 and neutral reviews around 17,000 by 2012.
- This trend may reflect greater consumer engagement, increased e-commerce adoption, or platform influence (e.g., Amazon's growth).
- The dominance of positive sentiment suggests brand perception is generally favorable, but businesses should remain alert to emerging negative sentiment trends as they scale.
- Monitoring sentiment shifts over time is vital for proactive customer experience management and reputation strategy.

Delivering to HULT INTERNATIONAL BUSINESS SCHOOL

All ▾ Exploratory Data Analysis

0 Team 3 Basket

Text and Sentiment Features: Sentiment Over Time

Seasonal Surge: Sentiment Peaks in Early Fall

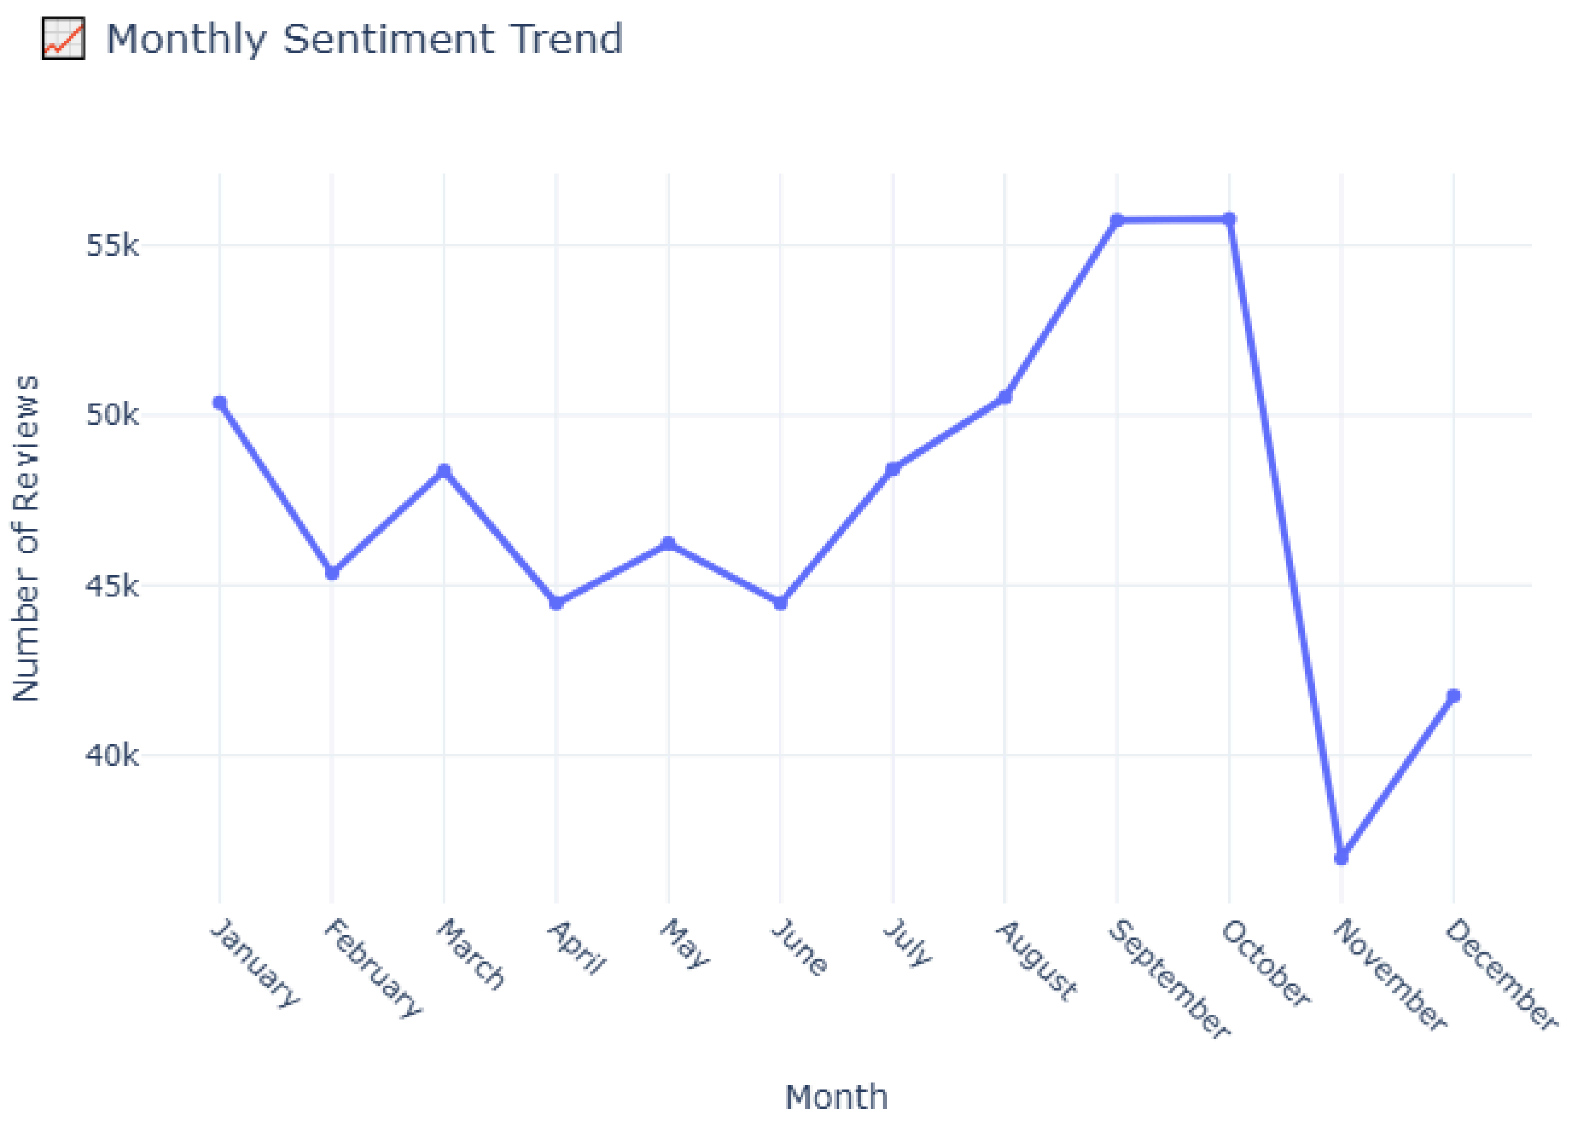

- Review activity rises steadily from June through October, peaking at over 55,000 reviews, indicating higher engagement during late summer to early fall.
- A sharp drop in November suggests a seasonal slowdown, possibly post-holiday sales or reduced user feedback before year-end.
- January starts strong, implying carryover from holiday shopping, while mid-year months show relatively stable, moderate activity.
- For businesses, this insight supports timing marketing campaigns and feedback requests around peak periods like August to October for maximum visibility.
- Aligning product launches or promotions with these peaks can boost customer interaction and review generation.

Exploratory Data Analysis: Correlation Heatmap

HULT
INTERNATIONAL
BUSINESS SCHOOL

Correlation Heatmap

Helpfulness Drives Data Relationships

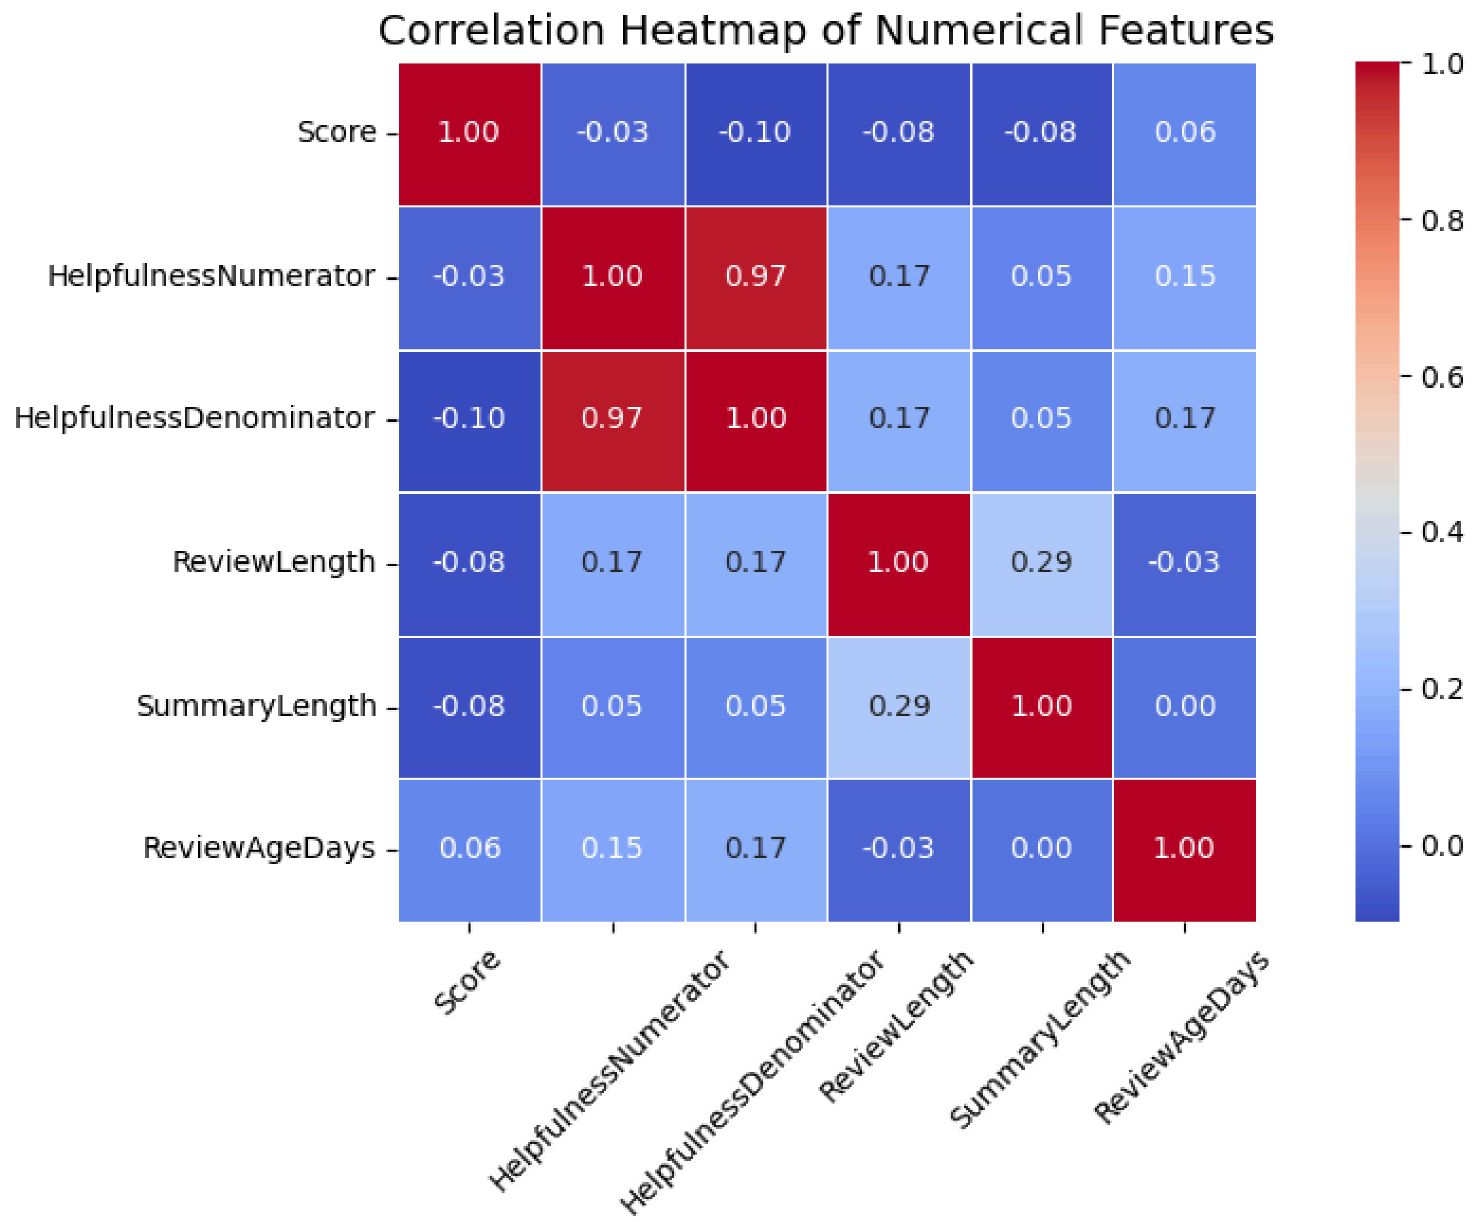

- HelpfulnessNumerator & HelpfulnessDenominator show an extremely strong positive correlation (0.97), confirming that users who rate reviews often do so in consistent proportions.
- ReviewLength moderately correlates with both helpfulness metrics (~0.17), suggesting longer reviews are slightly more likely to be rated helpful.
- SummaryLength has weak correlation with other features, indicating it adds little value for predictive modeling.
- Score is weakly negatively correlated with most variables (max |0.10|), suggesting star ratings are not strong indicators of review length or helpfulness votes.
- For business use, this matrix guides feature selection—prioritize helpfulness and review length over score when modeling customer feedback quality or engagement.

Exploratory Data Analysis: Time-Based Insights

HULT
INTERNATIONAL BUSINESS SCHOOL

Temporal Trends: Temporal Sentiment Trends by Product

Positive Buzz Dominates Top Products Over Time

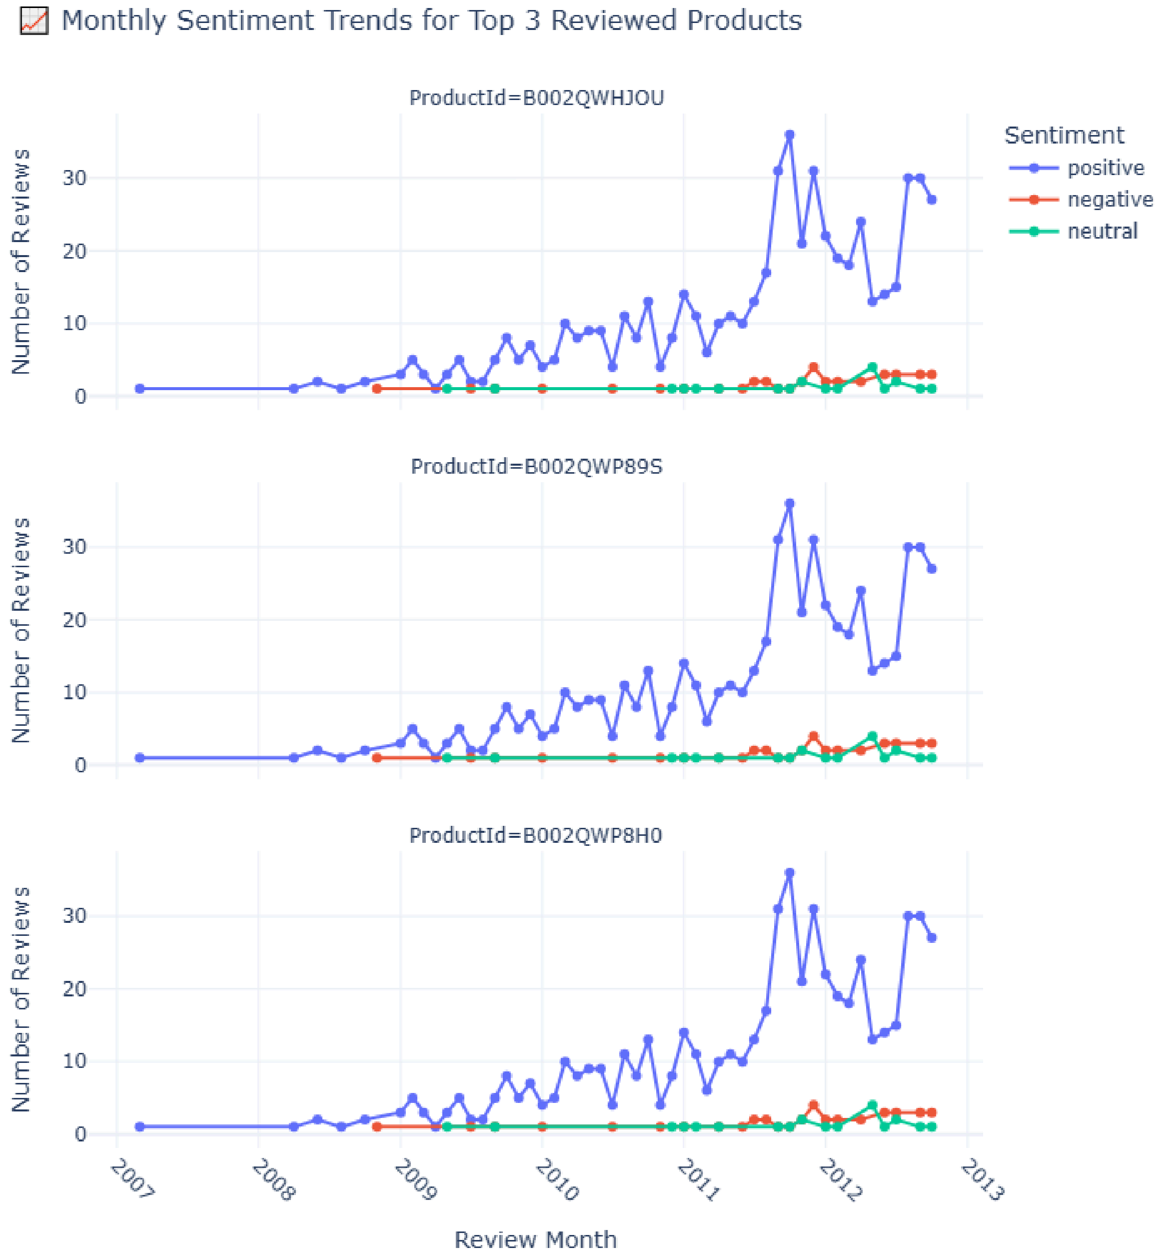

- All top 3 reviewed products show a steady increase in positive reviews, especially from 2010 onwards.
- Negative and neutral reviews remain minimal, suggesting strong customer satisfaction across high-volume items.
- The seasonal review spikes may indicate campaigns or restocking events, offering valuable cues for marketing timing.
- These trends reinforce product trustworthiness and brand loyalty, helping businesses prioritize which items to promote.

Temporal Trends: Sentiment Proportions Over Time by Product

Sentiment Share Holds Strong Despite Growing Negativity

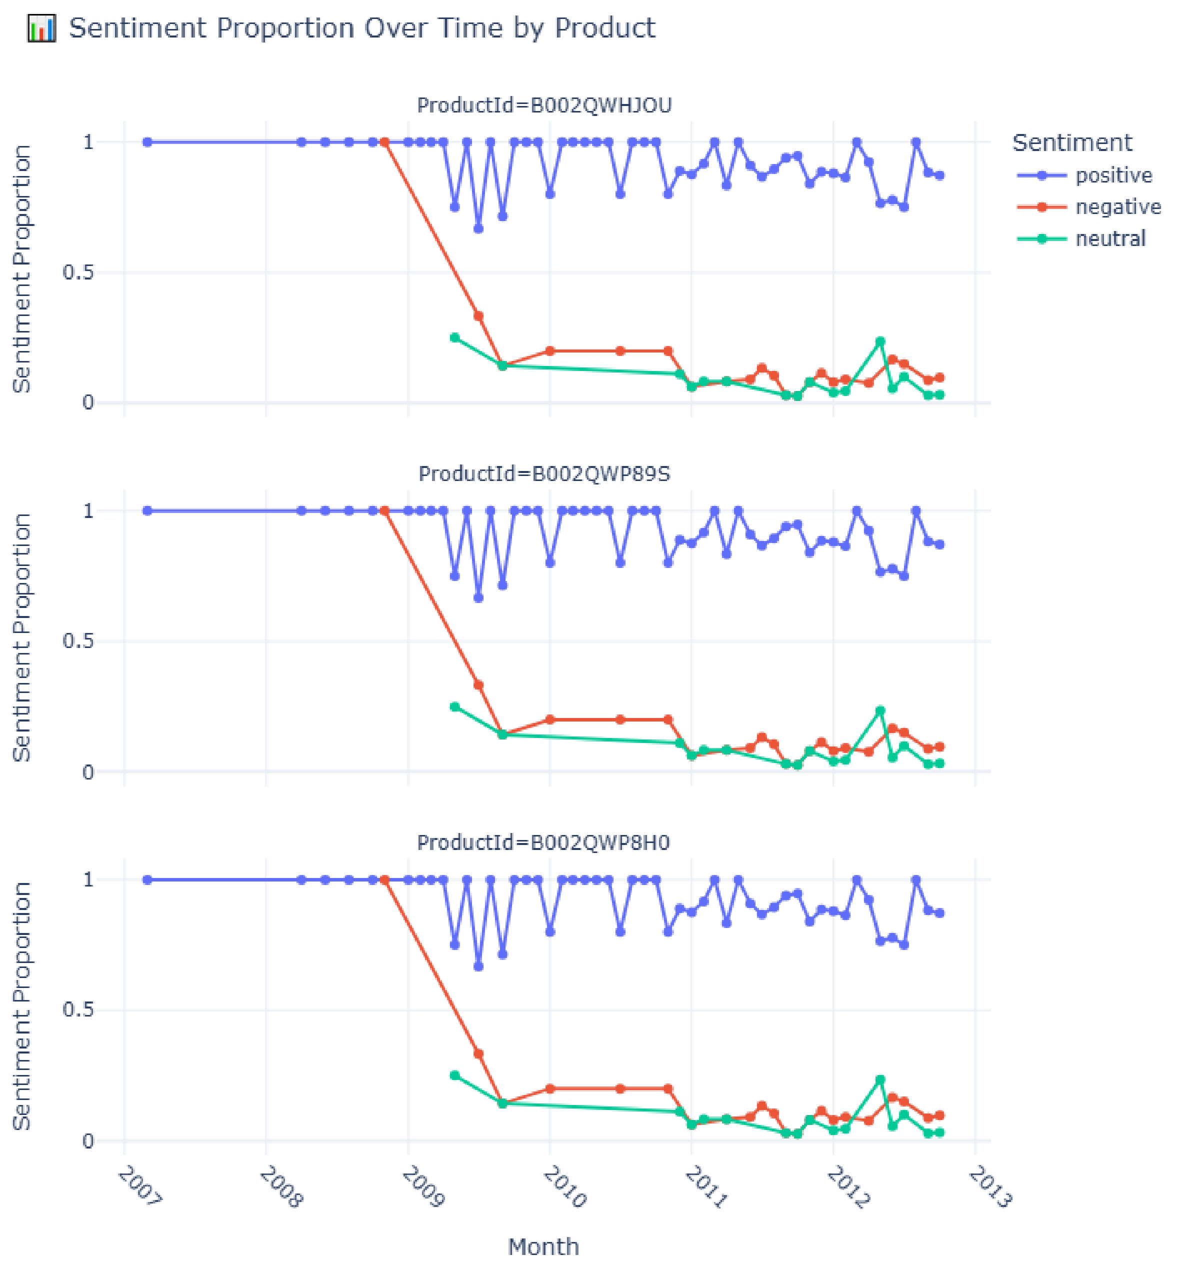

- Initially, all top 3 products had 100% positive sentiment, reflecting strong early reception.
- As review volume increased, negative and neutral proportions slightly rose, suggesting diverse user experiences over time.
- Despite this, positive sentiment consistently dominated, maintaining over 70% share across all periods.
- This pattern reveals resilient product perception, making them solid picks for continued promotion or feature.

Temporal Trends: Average Helpfulness Ratio Over Time by Product

Helpfulness Ratio Stays High with Occasional Dips

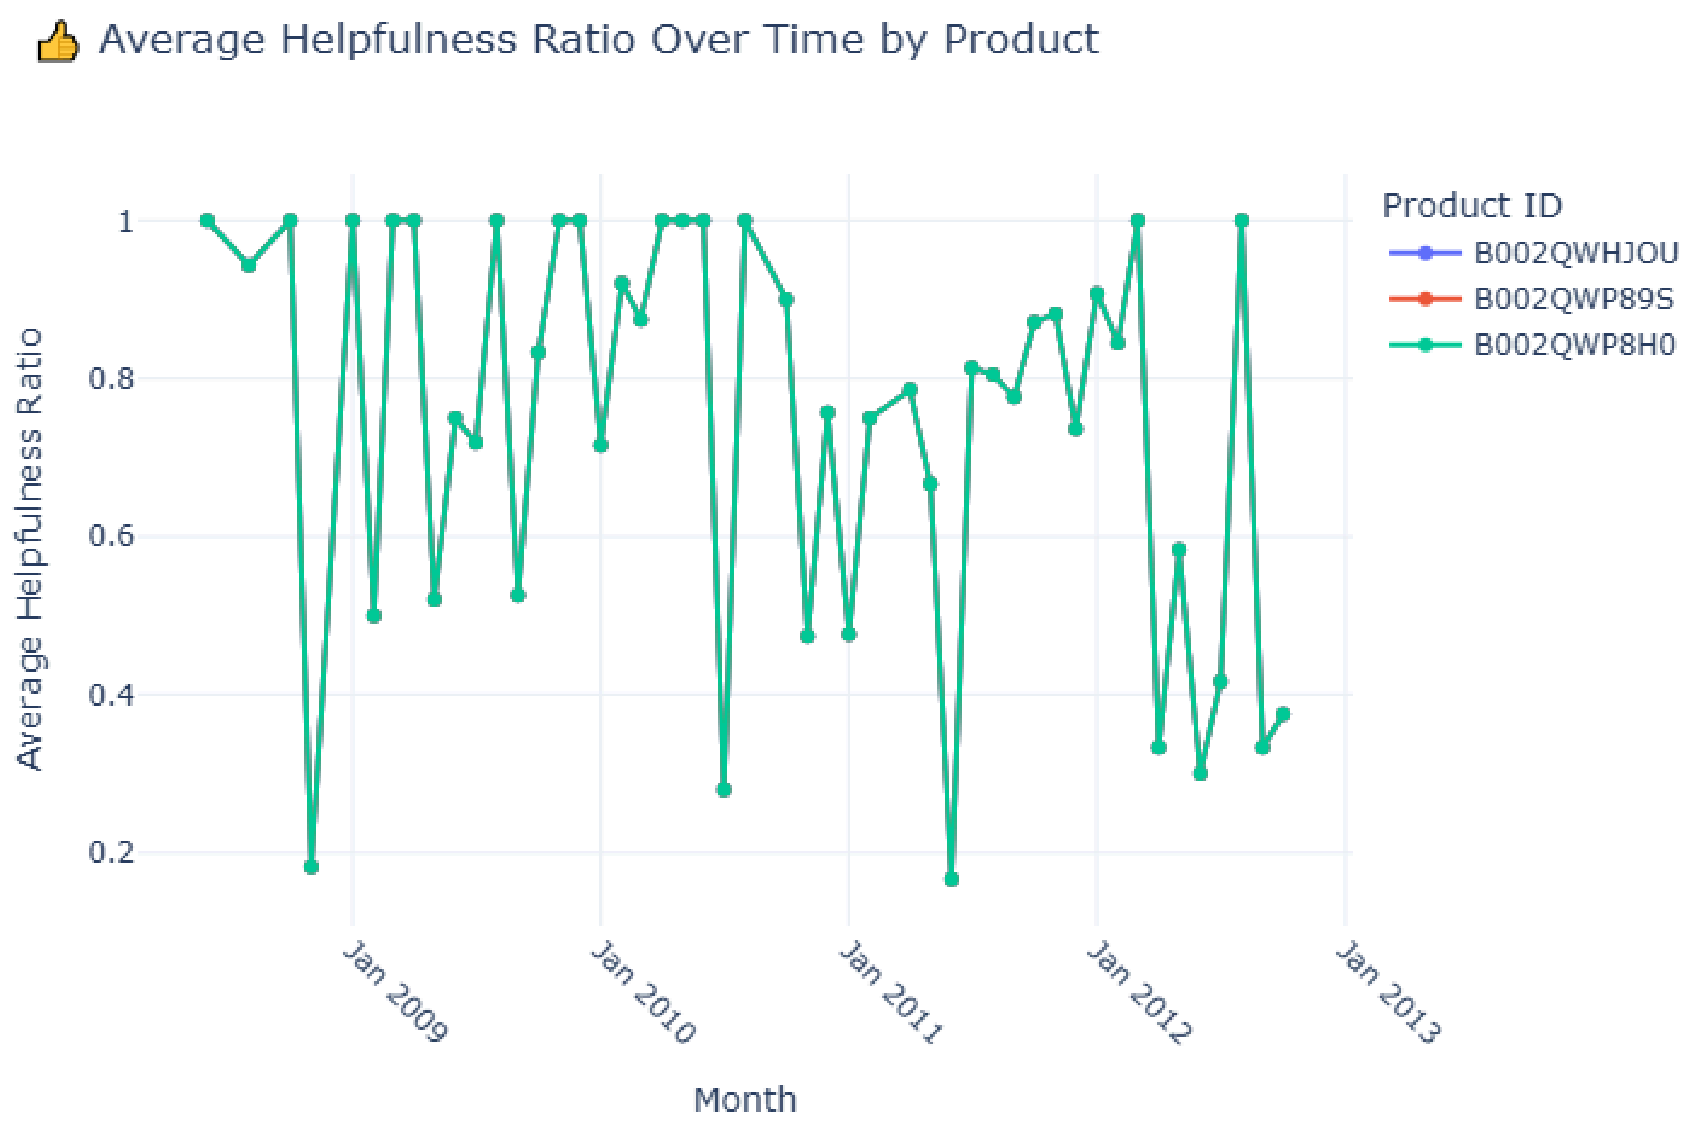

- Products maintain high average helpfulness ratios, often close to 1.0, signaling trusted, quality reviews.
- Intermittent sharp drops suggest influxes of low-rated or less informative feedback in certain months.
- These fluctuations highlight opportunities to improve review clarity or guide users on writing impactful reviews.
- Overall, strong helpfulness trends reinforce product credibility and user trust in peer recommendations.

Temporal Trends: Review Volume vs Helpfulness Ratio

High Volume Doesn't Always Mean High Helpfulness

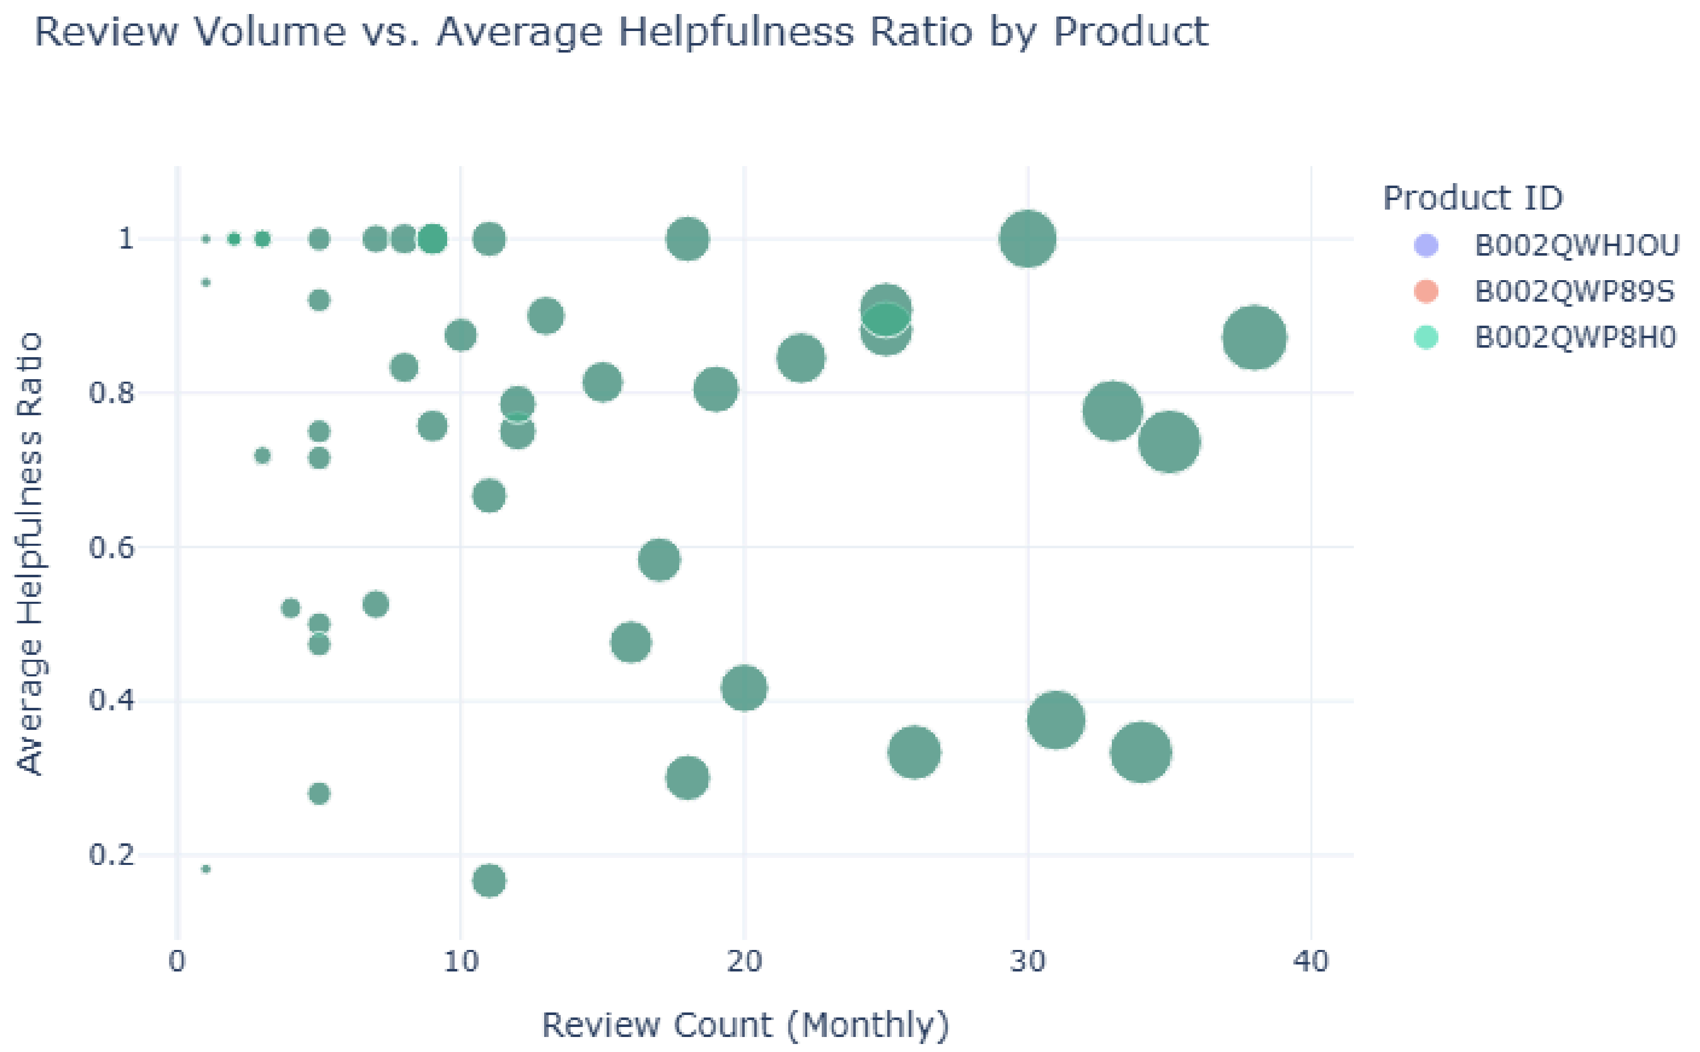

- Products with moderate review volumes (10–25/month) tend to show the highest helpfulness ratios.
- Extremely high volumes (30–40/month) sometimes correspond with lower helpfulness, possibly due to quality dilution or spam-like behavior.
- Sparse reviews can still be highly helpful, suggesting that thoughtful feedback matters more than quantity.
- Ideal strategy: Encourage fewer but higher-quality reviews to maximize perceived helpfulness.

Temporal Trends: Seasonal Sentiment Cycles by Product

Holiday Seasons Boost Positive Buzz

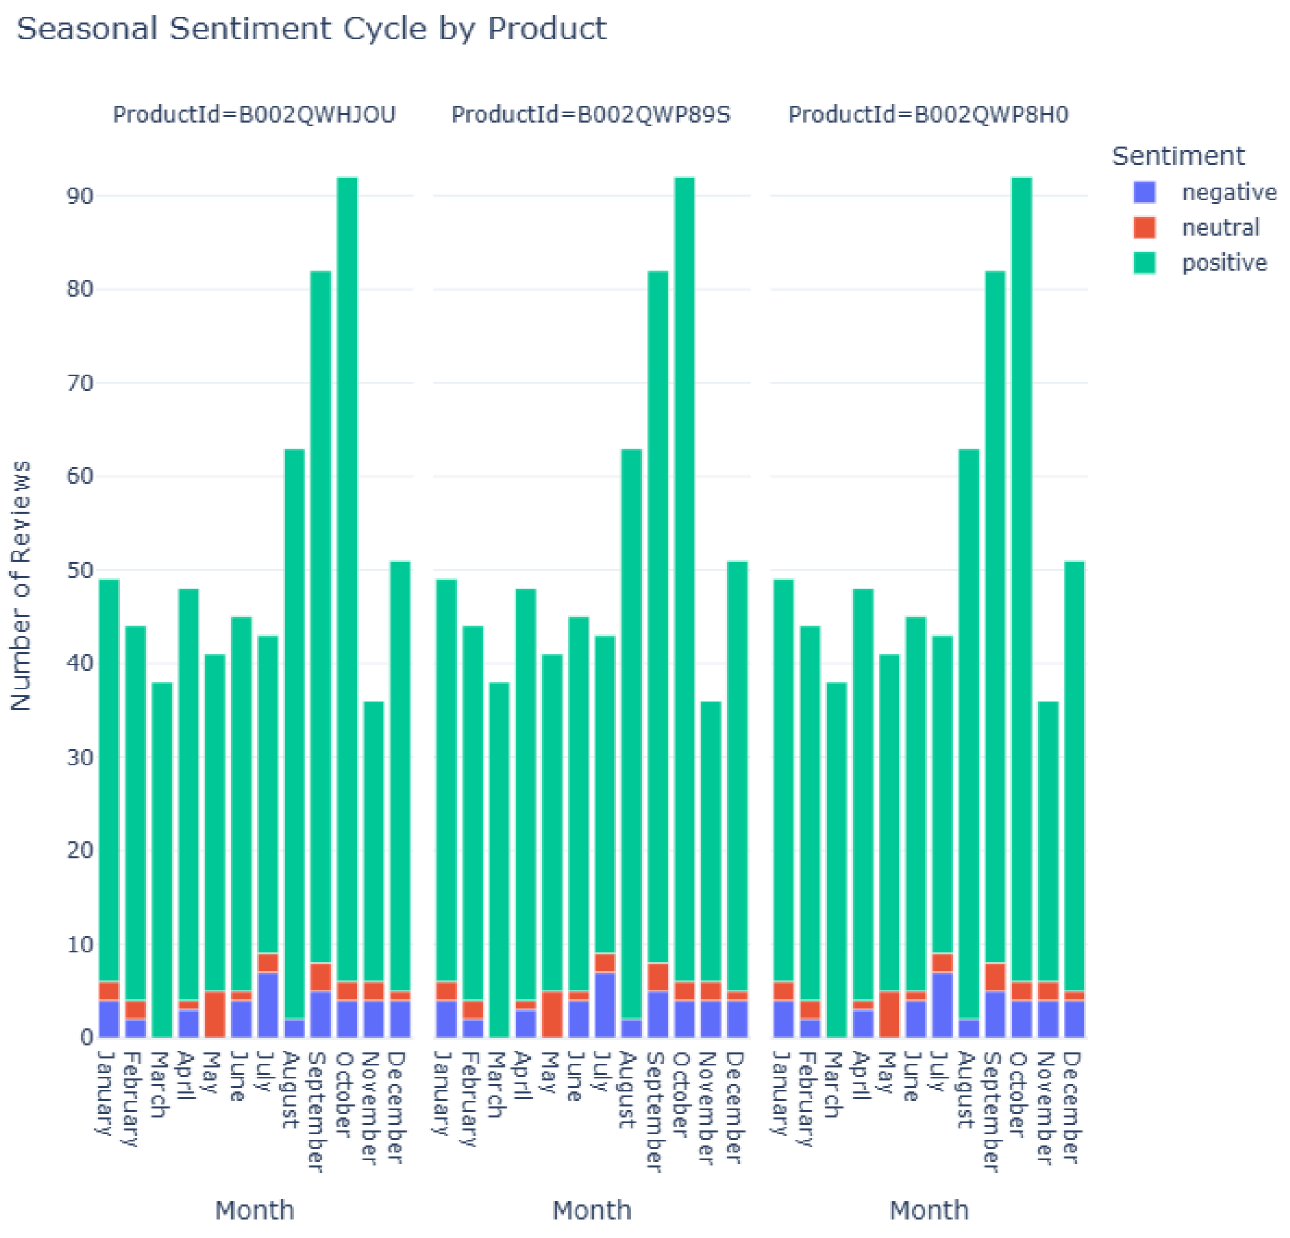

- All three products experience peaks in positive sentiment during Q4 (Oct–Dec), aligning with holiday shopping trends.
- December consistently shows the highest volume of positive reviews, suggesting gift-driven purchases.
- Negative and neutral sentiments remain relatively steady year-round, indicating low volatility in dissatisfaction.
- Marketers can leverage Q4 seasonality by timing campaigns, promotions, and influencer reviews to coincide with this surge.

Delivering to HULT INTERNATIONAL BUSINESS SCHOOL

All | Exploratory Data Analysis

0 Team 3 Basket

Temporal Trends: Compare Ratings Trends Alongside Sentiment

Consistent Ratings, Aligned Sentiments

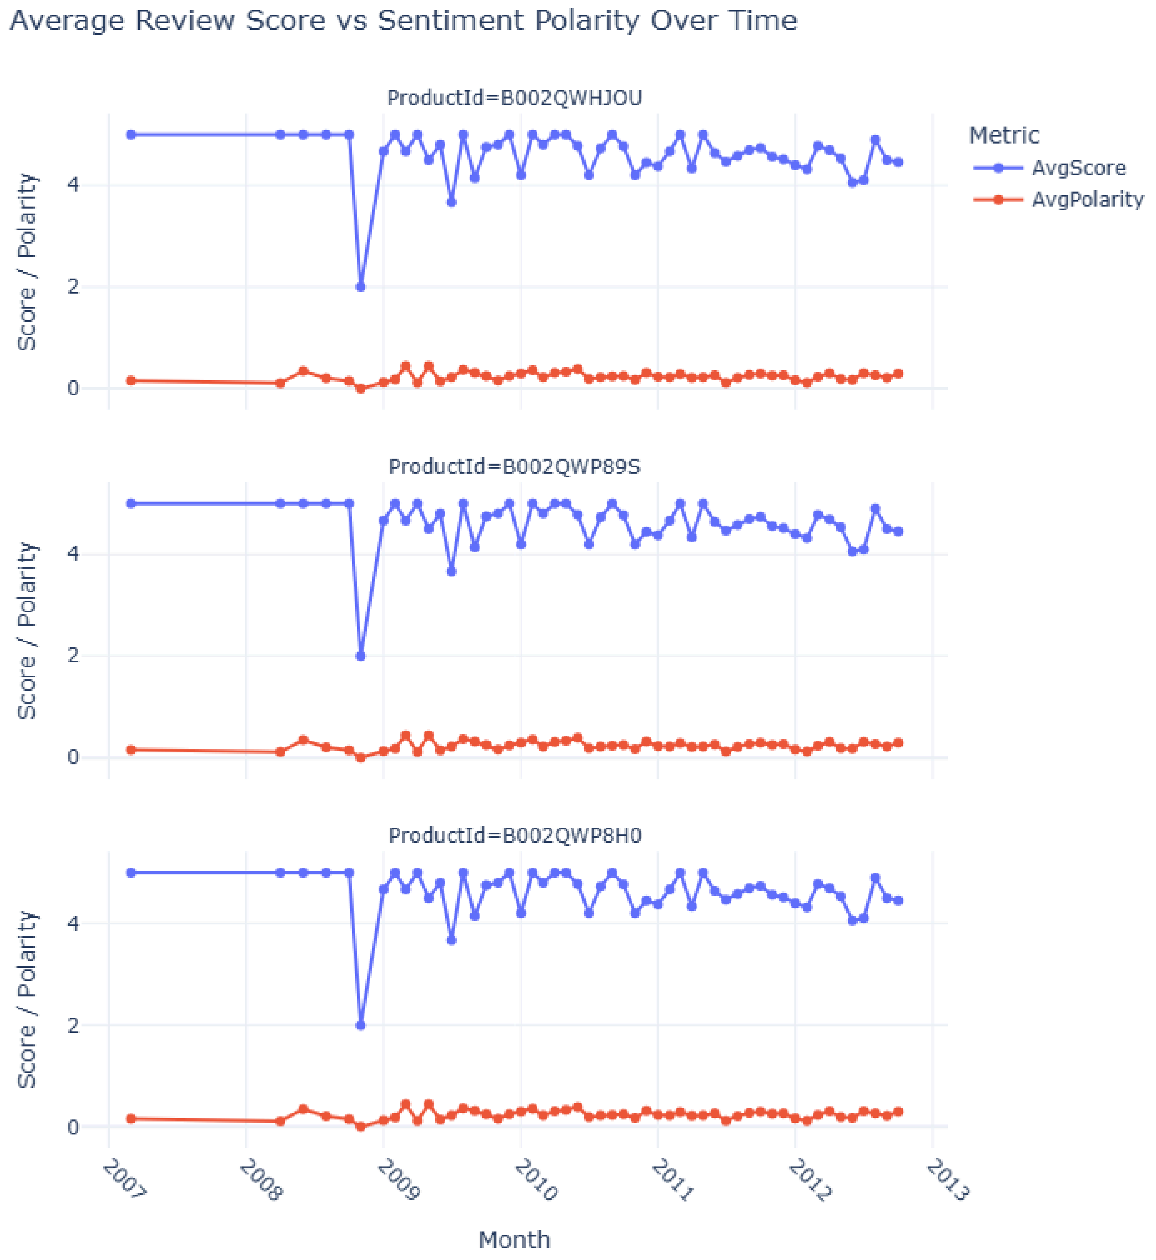

- All top products maintain a stable average review score (≈4.5), reflecting consistently high user satisfaction.
- Sentiment polarity trends closely mirror review scores, confirming the reliability of NLP sentiment as a proxy for customer feedback.
- The brief drop in 2008 is matched by a dip in polarity, indicating a true sentiment shift, not just a rating anomaly.
- This alignment validates the use of both star ratings and text sentiment for robust product performance monitoring.

Analytical Techniques: Sentiment Analysis

HULT
INTERNATIONAL BUSINESS SCHOOL

Sentiment Analysis

Amazon Shoppers Speak Loud: Positive Reviews Dominate Across the Years

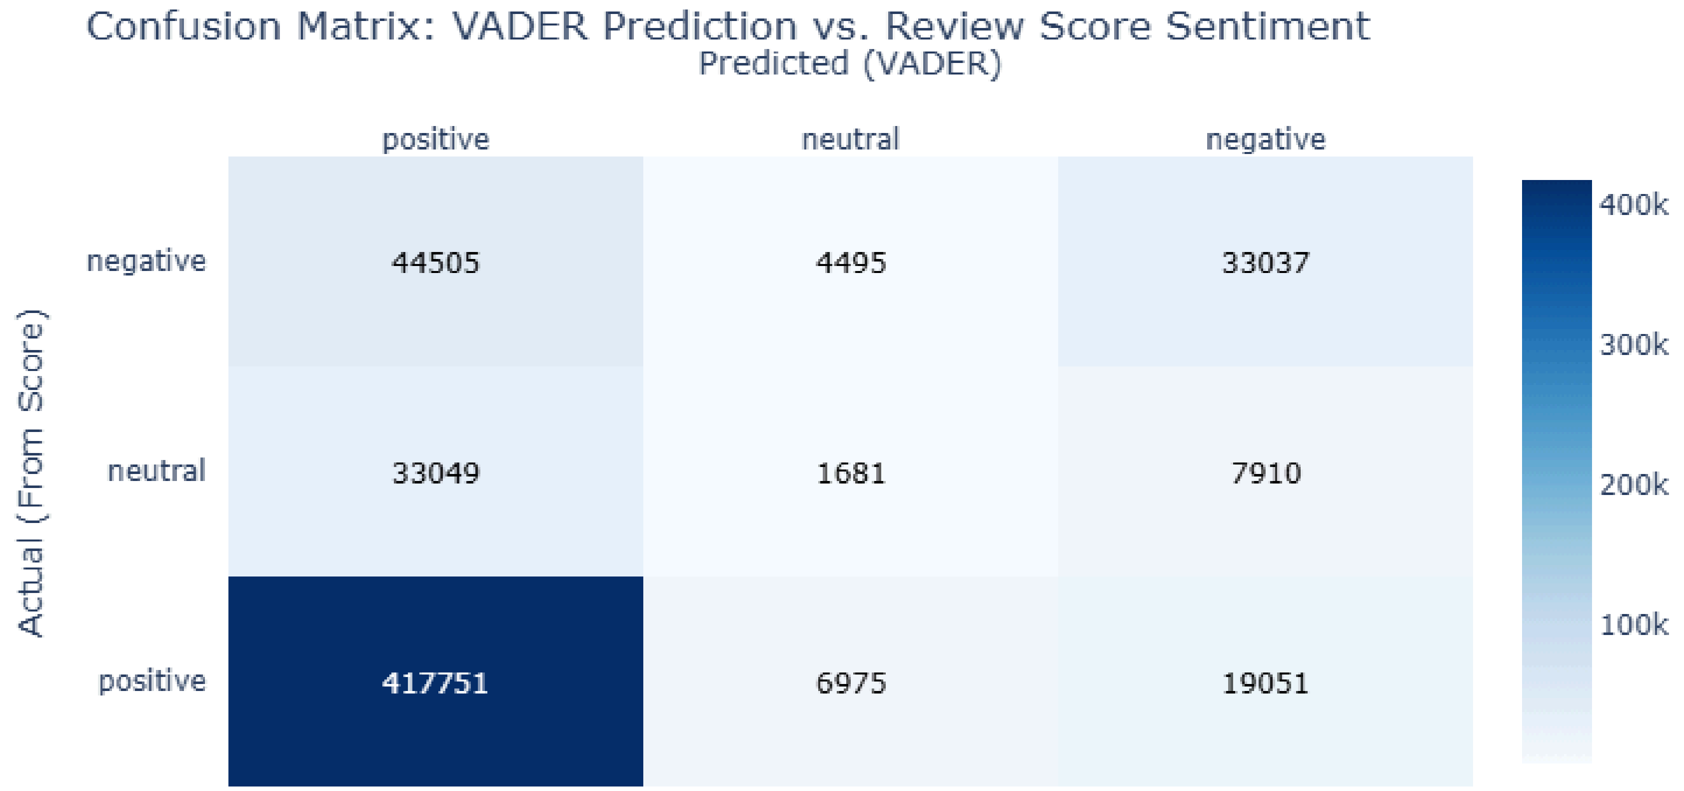

VADER Sentiment Classification Report:

	precision	recall	f1-score	support
negative	0.551	0.403	0.465	82037
neutral	0.128	0.039	0.060	42640
positive	0.843	0.941	0.890	443777
accuracy			0.796	568454
macro avg	0.507	0.461	0.472	568454
weighted avg	0.747	0.796	0.766	568454

- Overall Accuracy: 79.6% – VADER shows strong performance, especially for positive sentiments.
- Positive Sentiment:
  - Precision: 0.843 | Recall: 0.941 | F1-score: 0.890
  - Most correctly identified (417k+ true positives).
- Negative Sentiment:
  - Moderate performance with F1-score of 0.465
  - Significant misclassification as positive (33k+ cases).
- Neutral Sentiment:
  - Weak detection (F1-score: 0.060)
  - Often confused with negative or positive classes.
- Conclusion:
  - VADER excels at identifying positive reviews but struggles with neutral and nuanced tones.
  - Potential improvement with custom or ML-based classifiers.

Topic Modelling with Scikit-Learn LDA

What Customers Are Talking About: 5 Hidden Themes Uncovered

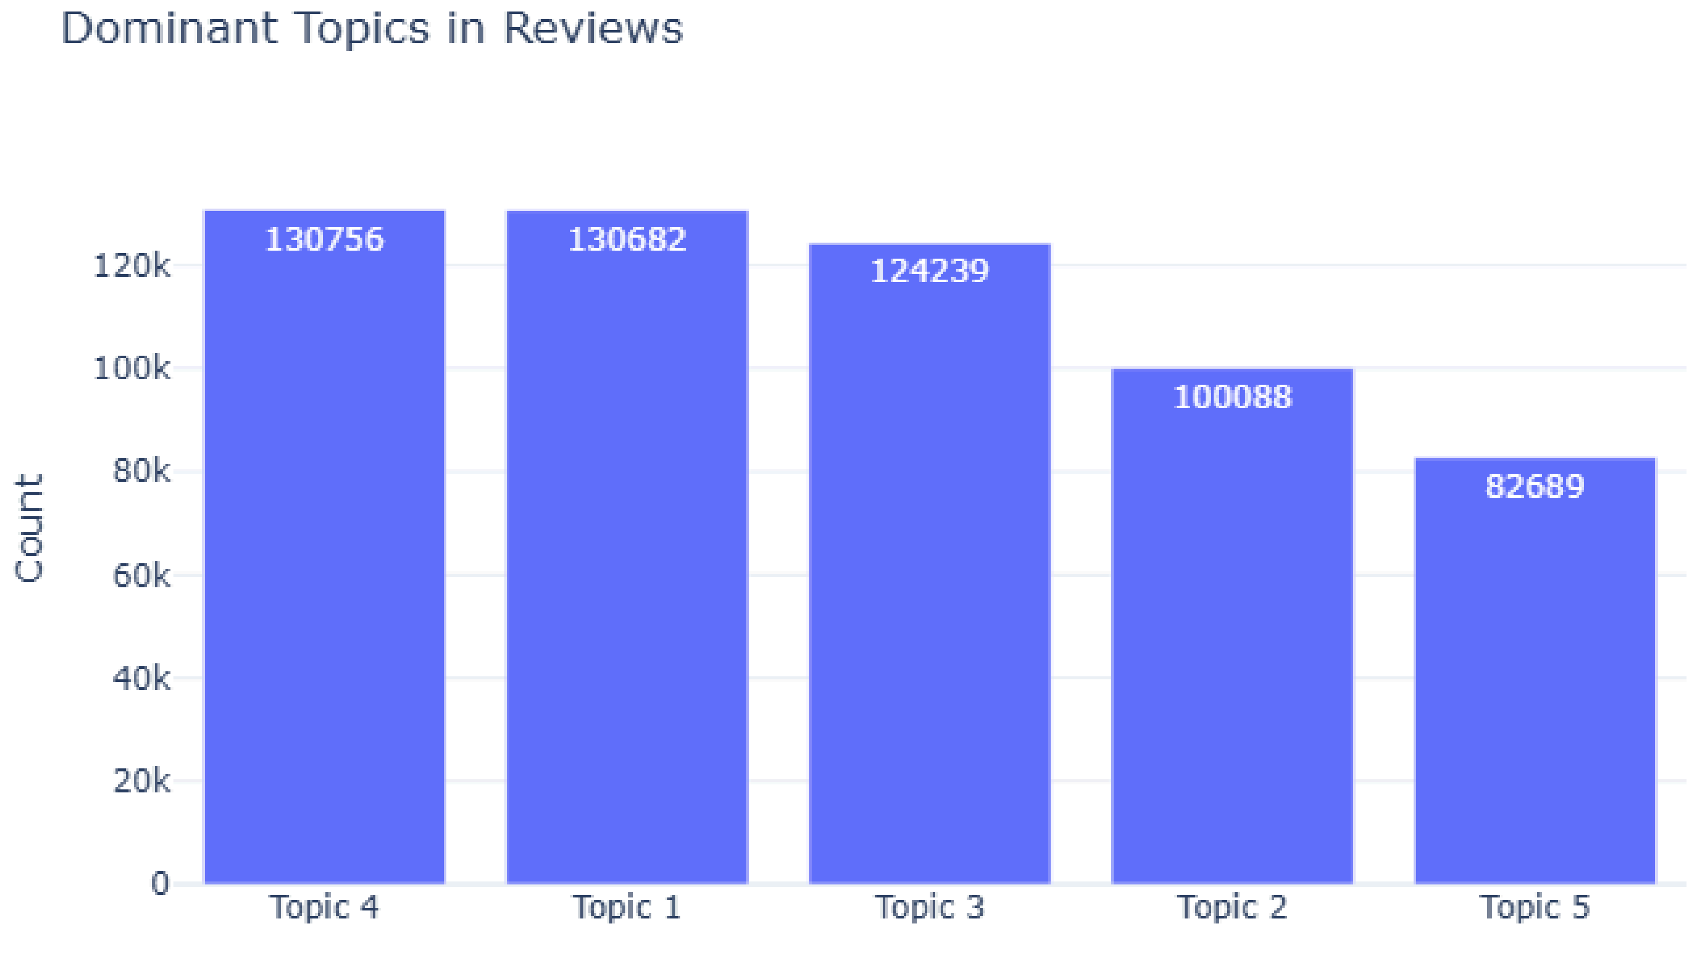

Topic 1: like, taste, good, chocolate, flavor, great, sugar, sweet, love, chips
Topic 2: food, dog, dogs, like, cat, treats, eat, loves, cats, love
Topic 3: coffee, tea, like, flavor, taste, cup, drink, good, water, green
Topic 4: great, amazon, price, product, good, store, buy, use, love, like
Topic 5: br, product, box, bag, dont, time, im, package, like, got

- Topic 1 – Sweet Cravings: Reviews center on taste, chocolate, and sugary snacks — a strong indicator of flavor satisfaction.
- Topic 2 – Pet Food Praise: Frequent mentions of dogs, cats, and treats show high engagement among pet owners.
- Topic 3 – Beverage Buzz: Words like coffee, tea, and green suggest loyal followings around popular drinks.
- Topic 4 – Value & Experience: Terms like “Amazon,” “price,” and “store” point to satisfaction with online purchases and product affordability.
- Topic 5 – Packaging & Delivery: Keywords highlight user experiences around boxes, bags, and timing—essential for logistics improvement.
- Insight: The LDA model reveals that product quality, flavor, pets, and pricing dominate user discussions, helping brands focus messaging and improvement efforts where they matter most.

Limitations & Future Work

HULT
INTERNATIONAL
BUSINESS SCHOOL

Limitations: Analysis

Data Bias:

Our dataset exhibited a significant bias towards positive reviews, with 78% of reviews being positive, 14% negative, and only 7.5% neutral. This distribution may not accurately reflect a balanced view of customer satisfaction.

Methodological limitations:

This study employed a specific set of analytical techniques, including VADER and LDA. However, other approaches may have different insights, and future research should consider exploring alternative methodologies.

Limited Scope:

The analysis focused on a specific product category, Amazon Fine Food Reviews, which may not be generalisable to other product categories or industries

Challenges in Sentiment Analysis:

The VADER sentiment classifier tends to overpredict positive sentiment, particularly in neutral or mixed reviews. This limitation may lead to an inaccurate representation of customer sentiment.

Limitations: Analysis

Data Quality Concerns	Insufficient Capture of Nuances
Limited Early Data: The availability of limited early data may distort the perceived value of specific feedback	The study´s relicanse on TF-IDF ectorisation and engineered features (e.g., review length, helpfulness ratio) may not fully capture the complexities and subtleties of customer feedback.
Review bias: Users with strong opinions may be more likely to post reviews, potentially leading to biased results	Biased results: The study´s findings may not accurately reflect customer sentiment and experiences
User participation limitations: many reviews remain unrated, which may impact the accuracy of our findings	Overlooked insights: More advanced text analysis techniques may uncover valuable insights that were not captured in this study

Future work

Exploring other product categories:

Analyse customer reviews from different product categories to identify common patterns and differences in sentiment and topic modelling.

Incorporating contextual information:

Integrate contextual factors, such as product price or marketing campaigns, to gain a more nuanced understanding of customer reviews and sentiment.

Using alternative analytical techniques:

Investigate the use of other analytical methods, such as deep learning models or aspect-based sentiment analysis, to improve the accuracy and insights gained from customer reviews

Developing a more comprehensive review quality assessment:

Develop a framework to assess review quality and identify potential biases or fake reviews.

Proposed Enhancements

HULT
INTERNATIONAL
BUSINESS SCHOOL

Proposed Enhancements for Amazon's Review Ecosystem (2025 Edition)

1. AI-Powered Summarization of Product Sentiment

- Deploy AI-generated summary cards at the top of product pages, displaying:
  Top Positives (e.g., "delicious flavor," "great packaging")
  Common Complaints (e.g., "received expired product," "packaging leaks")
- Use NLP to auto-generate pros/cons summaries from recent reviews.

2. Contextual Review Highlighting Based on User Intent

- Let users filter and view reviews by topic category (e.g., "delivery experience" or "taste profile") and usage intent (e.g., "keto," "gifting," "kids").

3. Sentiment-Triggered Seller Feedback Loop

- Introduce an AI-powered alert system that:
- Flags products with sudden dips in sentiment score (e.g., 4.5 → 3.2 in <3 months).
- Notifies third-party sellers for quality control.

Proposed Enhancements for Amazon's Review Ecosystem (2025 Edition)

4. Reviewer Network Insights for Personalization

- Build a "People like you also liked…" system:
- Recommend products reviewed positively by users in your taste/preference cluster.
- Let shoppers follow or save reviewer profiles based on food preferences.

5. Bias Mitigation & Authenticity Verification

- Insight: Heavily skewed reviews: ~75% are 5-star, only ~5% are 1-star — suggesting potential bias or selection effect. Enhancements:
- Show a balanced sentiment distribution chart.
- Highlight "critical but helpful" reviews to counteract review inflation.
- Introduce Verified Usage Tags (e.g., "Consumed within 30 days," "Used in baking").

6. Enhance Review Trustworthiness

- Introduce Verified Purchase Badge for Reviews
- Encourage authentic feedback and reduce fake/biased reviews.
- Weight Reviews by Helpfulness & User Credibility
- Rank and prioritize reviews based on user history and helpfulness score.

Proposed Enhancements for Amazon's Review Ecosystem (2025 Edition)

7. Community Engagement

- Gamify Helpful Review Contributions. Introduce badges, levels, or incentives for top reviewers in each category.
- Enable Reviewer Follow Options. Let users follow trusted reviewers for personalized recommendations.

8. Voice-Generated Summaries of Reviews

- Introduce audio summaries or Alexa-style voice playback of key review insights (e.g., “Most users love this for smoothies. Some complained about late delivery.”)
- Justification: 2025 is voice-first — many users scroll less and listen more. Audio summaries add accessibility and reduce cognitive load during shopping.

9. Sustainability Sentiment Tagging

- Allow users to filter reviews that mention eco-friendliness, plastic usage, organic, etc.
- This matters to Eco-conscious people and Gen Z, who often check reviews for eco claims.

10. AI Auto-Tags Based on Context

- Detect and auto-tag when and how the product is used: “Great for school lunches”, “Used in protein shakes”, “Perfect for camping”
- Detect keywords related to dietary restrictions or allergens and auto-warn users. “Not really gluten-free” or “Caused nut allergy reaction”
- Flag these for visibility or moderation

Summary

- The static star ratings and unstructured reviews no longer meet evolving customer expectations. These enhancements can:
- Cut users product research time by 35–50%
- Improve seller accountability
- Increase review helpfulness and personalization
- Enable proactive issue detection via sentiment trends

Implication of Insights to the Chosen Business/Industry

Amazon-Specific Implifications

Integrate Review Sentiment into Forecasting and Inventory Decisions

Amazon can embed review sentiment analysis into its demand forecasting models. A sustained rise in positive sentiment for certain product categories could signal upcoming surges in demand, helping optimize inventory allocation, reduce stockouts, and minimize holding costs.

Monitor Reviewer Network Clusters for Emerging Trends

By applying social network analysis to its reviewer base, Amazon can identify micro-communities and trendsetters. These reviewer clusters often reflect niche audiences whose behaviors can predict broader consumer shifts, providing early signals for emerging products or preferences.

Incorporate Longitudinal Sentiment Analysis for Seller Scorecards

Instead of only snapshot metrics, Amazon could assess sentiment change over time for each seller. This trajectory-based insight enables a more robust evaluation of seller performance, buyer satisfaction, and brand trajectory.

Develop Sentiment-Driven Dynamic Pricing Models

Amazon can experiment with real-time pricing strategies that consider product sentiment. Positive sentiment momentum may justify a price premium, while negative shifts can trigger automated discounts, all without relying on historical sales data alone.

Enable Sentiment-Indexed Advertising Optimization

Amazon Ads can optimize placement and spend by aligning sponsored listings with sentiment-rich reviews. Ads for products with positive sentiment patterns can be prioritized in placements, while those with neutral or declining sentiment can be paused or retargeted with updated messaging.

Industry-Specific Implifications

Standardize Cross-Platform Review Data Protocols

As consumers increasingly shop across marketplaces, the industry must establish interoperable review data standards. This would allow for portable reputation systems, enhancing trust across platforms.

Introduce Review Sentiment as a KPI in ESG Reporting

Brands can integrate consumer sentiment metrics into Environmental, Social, and Governance (ESG) dashboards, particularly under the "S" (social impact) pillar. Sentiment trends can indicate consumer perception of ethical sourcing, diversity, and sustainability efforts.

Enable Cross-Industry Benchmarking Through Sentiment Intelligence

E-commerce platforms and market research firms can collaborate to create anonymized, aggregated sentiment benchmarks by category and region. This enables strategic decision-making on product innovation, brand health, and regional expansion based on collective consumer voice.

Foster Ethical AI Governance in Review Interpretation

Given the rise of AI-generated summaries and sentiment tagging, industry actors must establish ethical guidelines and transparency standards for how algorithms interpret and surface user feedback, ensuring fairness, bias mitigation, and consumer trust.

Institutionalize Review Sentiment in Brand Equity Valuation Models

For brand managers and investors, review sentiment analysis should become a formal variable in brand valuation methodologies, signaling reputation risk, loyalty potential, and long-term equity resilience in a way traditional survey-based methods can’t match.

Conclusion

BUSINESS PROBLEM VS EVIDENCE FROM THE DATA VS OBJECTIVE

BUSINESS PROBLEM	EVIDENCE FROM THE DATA	OBJECTIVE
01 What makes a review helpful or influential?	Longer reviews had higher helpfulness ratios. Helpfulness is also bimodal (either ignored or fully endorsed), highlighting engagement patterns.	Trends in review helpfulness
02 How does customer sentiment shift over time and across products?	Temporal sentiment analysis showed rising positive trends over time (esp. after 2006), with seasonal review surges in Q3–Q4. Product-specific trends confirmed sentiment consistency or gradual decline.	Sentiment dynamics over time and products
03 Which product aspects drive positive or negative feedback?	Topic modeling (LDA) extracted 5 main themes: sweet/snack flavor, pet food, beverages, value/experience, and packaging/delivery. These revealed what customers focus on most.	Thematic topics in feedback

RECOMMENDED ACTION

INSIGHT	IMPLICATION	RECOMMENDED ACTION
Longer reviews = more helpful	Review quality matters more than volume	Prompt users to write structured, detailed reviews
Positive sentiment dominates	May mask dissatisfaction or subtle issues	Use custom ML tools to detect nuanced sentiment
Review helpfulness is bimodal (0 or 1)	Many reviews go unrated or are ignored entirely	Improve UI design to highlight useful reviews and prompt helpful votes
Topic modeling revealed 5 consistent themes	Flavor, pet care, beverages, value, and delivery drive feedback focus	Align product marketing and R&D with these core customer concerns
Seasonal sentiment trends observed (Q3–Q4 peaks)	Customer engagement is time-sensitive	Schedule marketing campaigns and review requests during peak seasons
Top users and products dominate activity	A small group heavily influences review patterns	Monitor power users and high-volume products for quality control
Star ratings are not strong predictors alone	Text analysis adds needed depth beyond rating averages	Combine sentiment, length, and helpfulness for product evaluation

BUSINESS STRATEGY

01 PRODUCT DEVELOPMENT STRATEGY (Prioritize Based on Voice of the Customer)

- Insight: Topic modeling identified key focus areas: taste, pet food, beverages, value perception, and delivery experience.
- Strategy:
  - Use review text themes to guide product improvement priorities (e.g., adjust flavor profiles, improve packaging).
  - Develop feedback loops between customer sentiment and product design teams.
- Outcome Target: Minimize repeat complaints and accelerate feature innovation based on direct user feedback.

02 MARKETING STRATEGY (Align Messaging with Sentiment & Seasonality)

- Insight: Reviews are overwhelmingly positive, especially in Q3–Q4.
- Strategy:
  - Launch campaigns during seasonal sentiment peaks (August–December).
  - Highlight long, helpful positive reviews in product pages or ads to enhance trust.
  - Personalize content based on top-reviewed product themes (e.g., “Trusted by pet lovers” or “Flavor people rave about”).
- Outcome Target: Increase review volume and product conversion by timing and tailoring marketing efforts.

BUSINESS STRATEGY

03 CUSTOMER EXPERIENCE STRATEGY (Elevate Review Quality and Trust Signals)

- Insight: Long, well-written reviews are more helpful and influential. Many reviews go unrated or unseen.
- Strategy:
  - Encourage structured reviews with smart prompts during checkout follow-up (e.g., “Tell us about taste, delivery, and value”).
  - Redesign the UI for helpfulness voting, to make quality reviews more visible.
  - Implement helpfulness scoring + sentiment analysis for review ranking.
- Outcome Target: Build a more trustworthy review ecosystem that guides new buyers and improves brand credibility.

Team 3

Angel LANTO
Fransiscus XAVERIUS
Hedda PEDERSEN
Mansi LODHIA
Rafael Martínez Quiles
Zahirah SMITH